RISC-V Pipeline Hazards!

Instructor: Sean Farhat

Great Idea #4: Parallelism

Software   *Hardware*

- **Parallel Requests**

  Assigned to computer

  e.g. search "Garcia"

- **Parallel Threads**

  Assigned to core

  e.g. lookup, ads

- **Parallel Instructions**

  > 1 instruction @ one time

  e.g. 5 pipelined instructions

- **Parallel Data**

  > 1 data item @ one time

  e.g. add of 4 pairs of words

- **Hardware descriptions**

  All gates functioning in
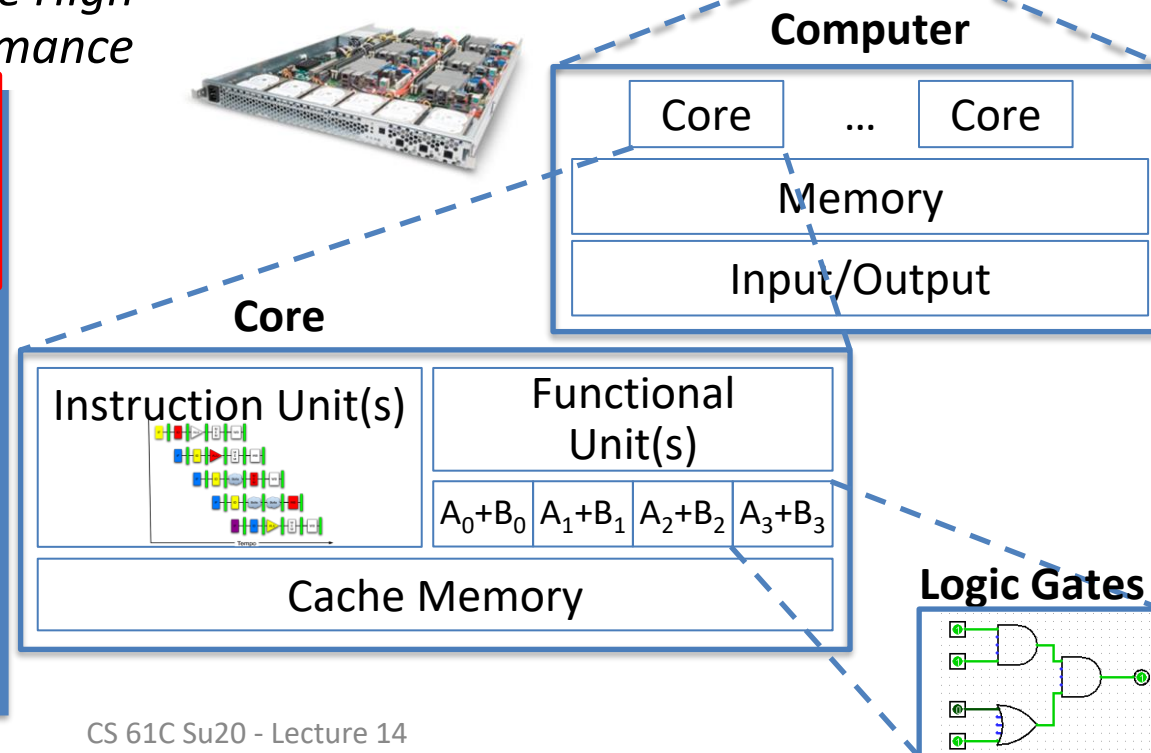 parallel at same time

Leverage Parallelism & Achieve High Performance

Warehouse Scale Computer

Smart Phone

Computer

| Core | ... | Core |

Memory

Input/Output

Core

Instruction Unit(s)

Functional Unit(s)

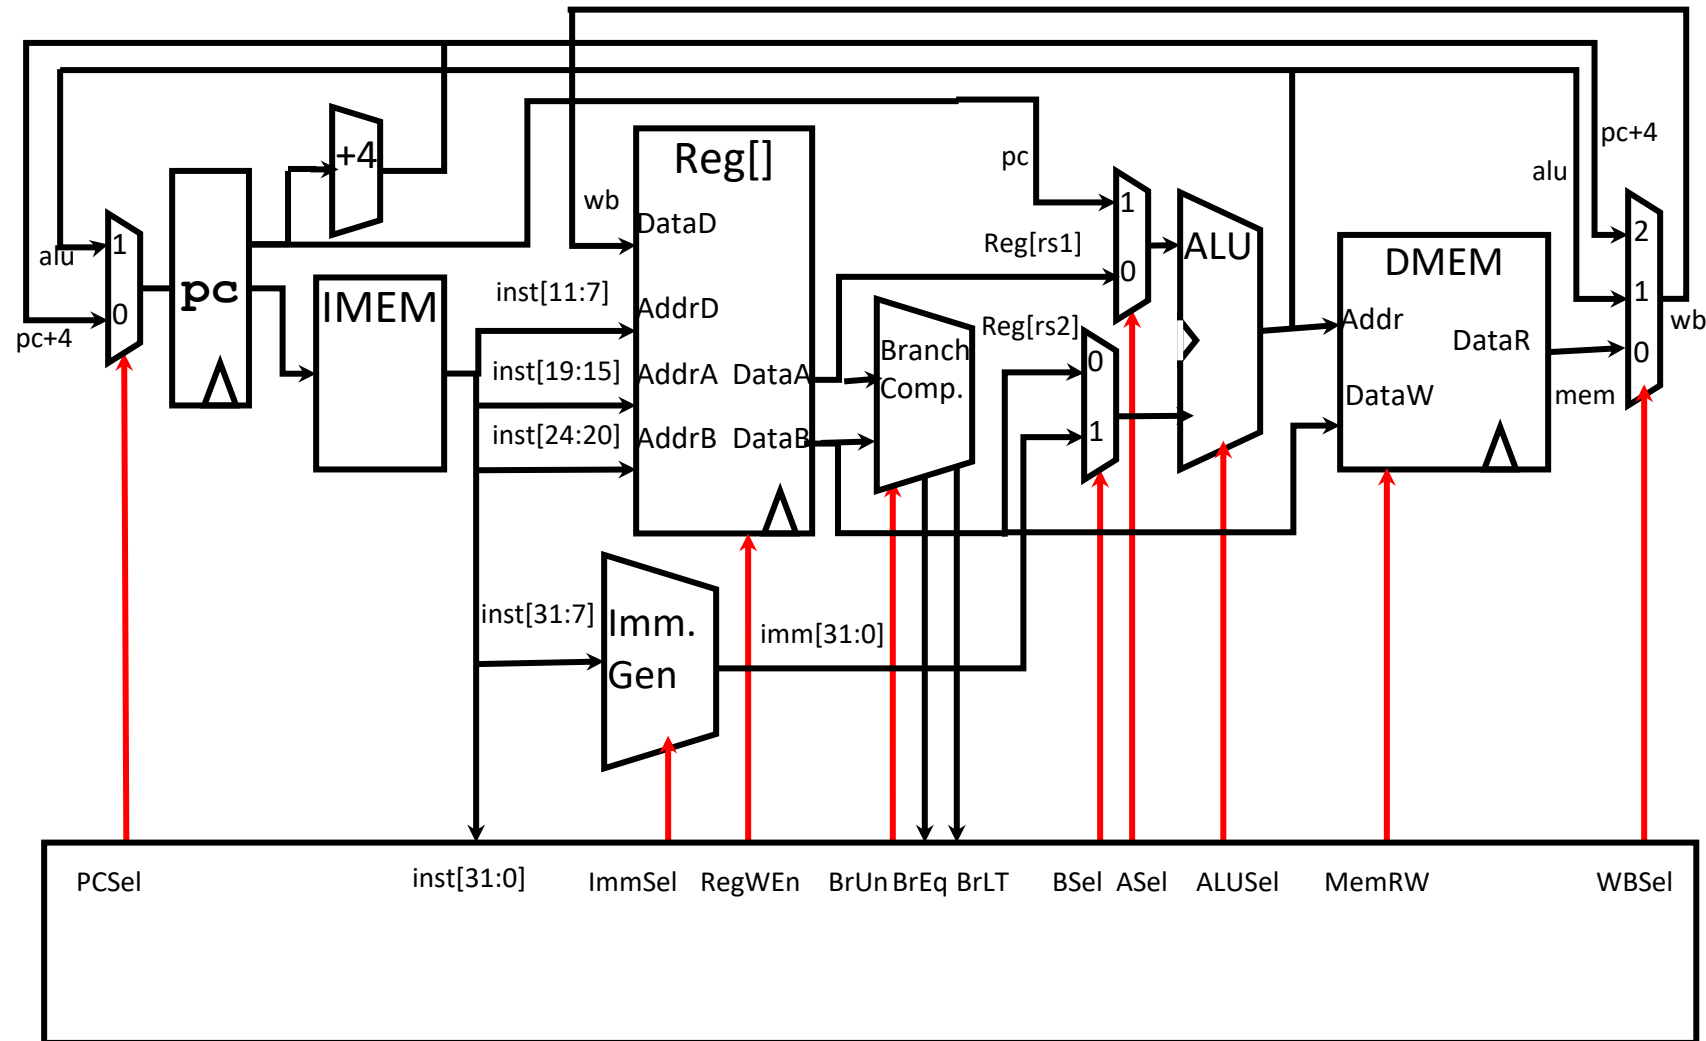| A_0+B_0 | A_1+B_1 | A_2+B_2 | A_3+B_3 |

Cache Memory

Logic Gates

Review of Last Lecture

- Implementing controller for your datapath
  - Take decoded signals from instruction and generate control signals
- <mark>Pipelining improves performance by exploiting Instruction Level Parallelism</mark>
  - <mark>5-stage pipeline for RISC-V: IF, ID, EX, MEM, WB</mark>
  - Executes multiple instructions in parallel
  - Each instruction has the same latency
  - What can go wrong???
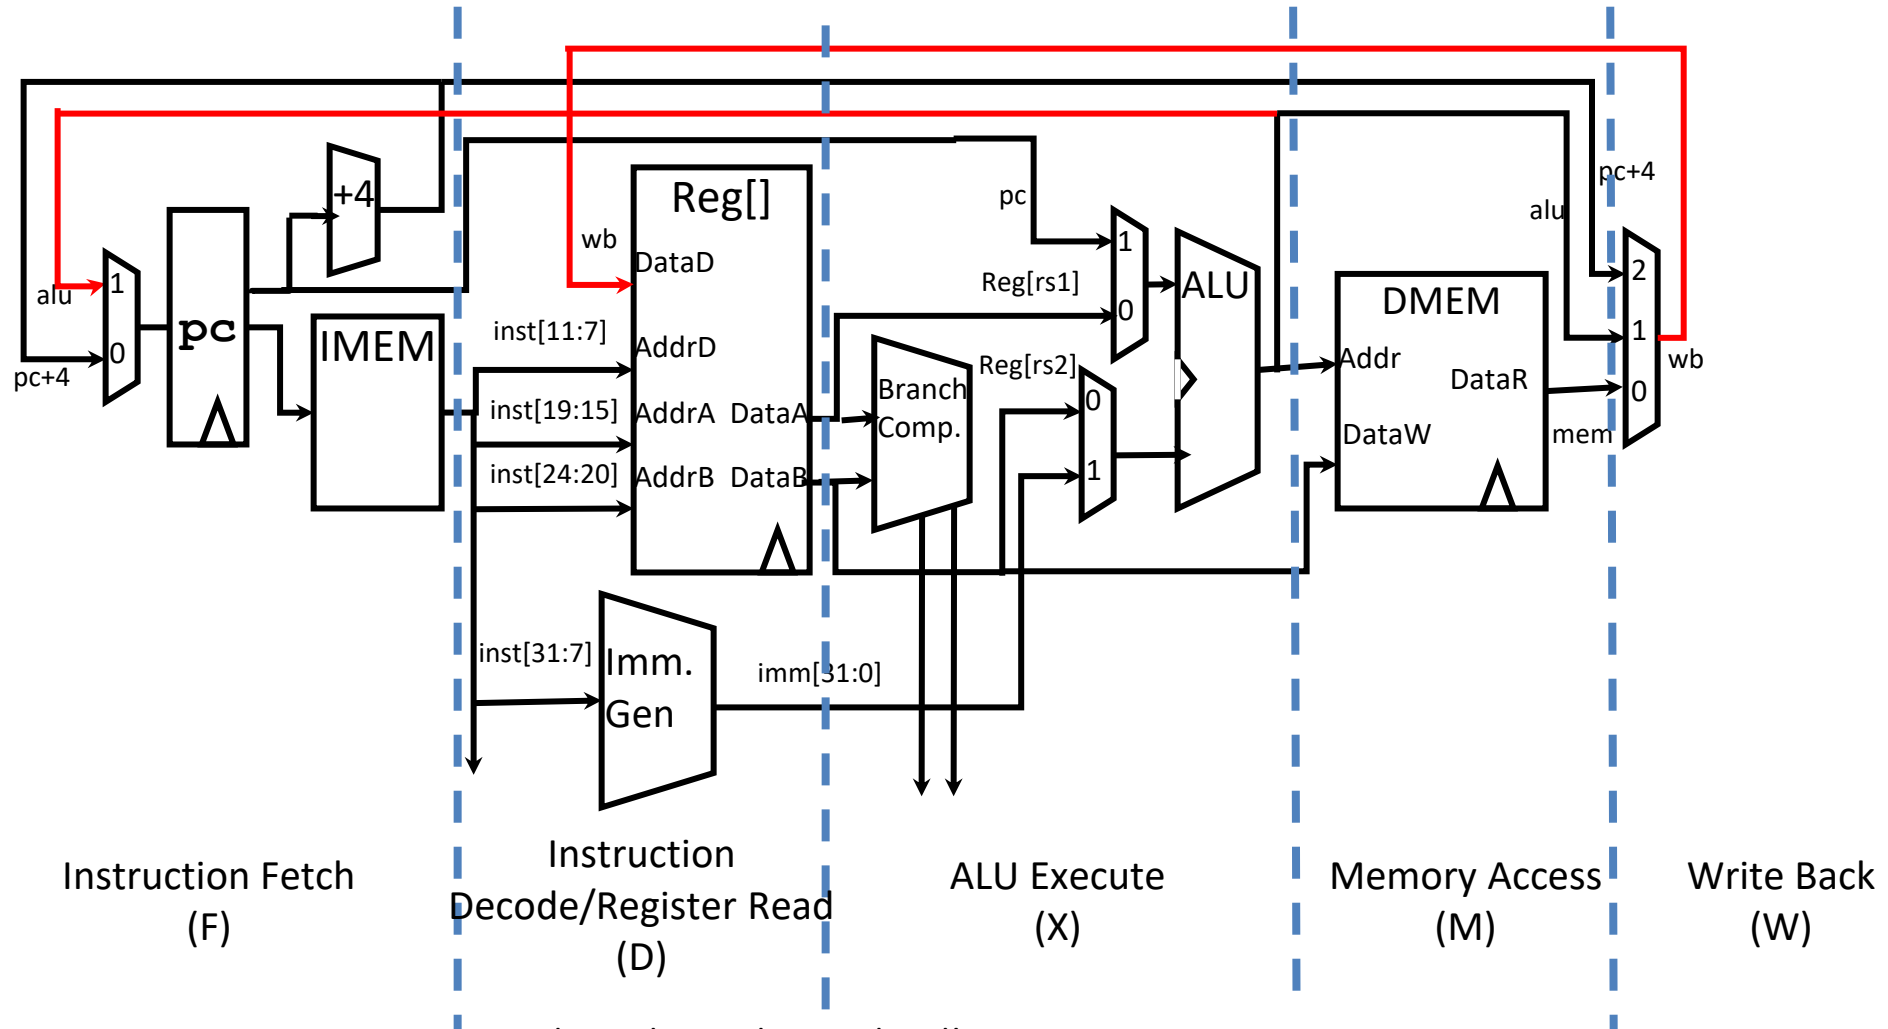
Agenda

- RISC-V Pipeline
- Hazards
  - Structural
  - Data
    - R-type instructions
    - Load
  - Control
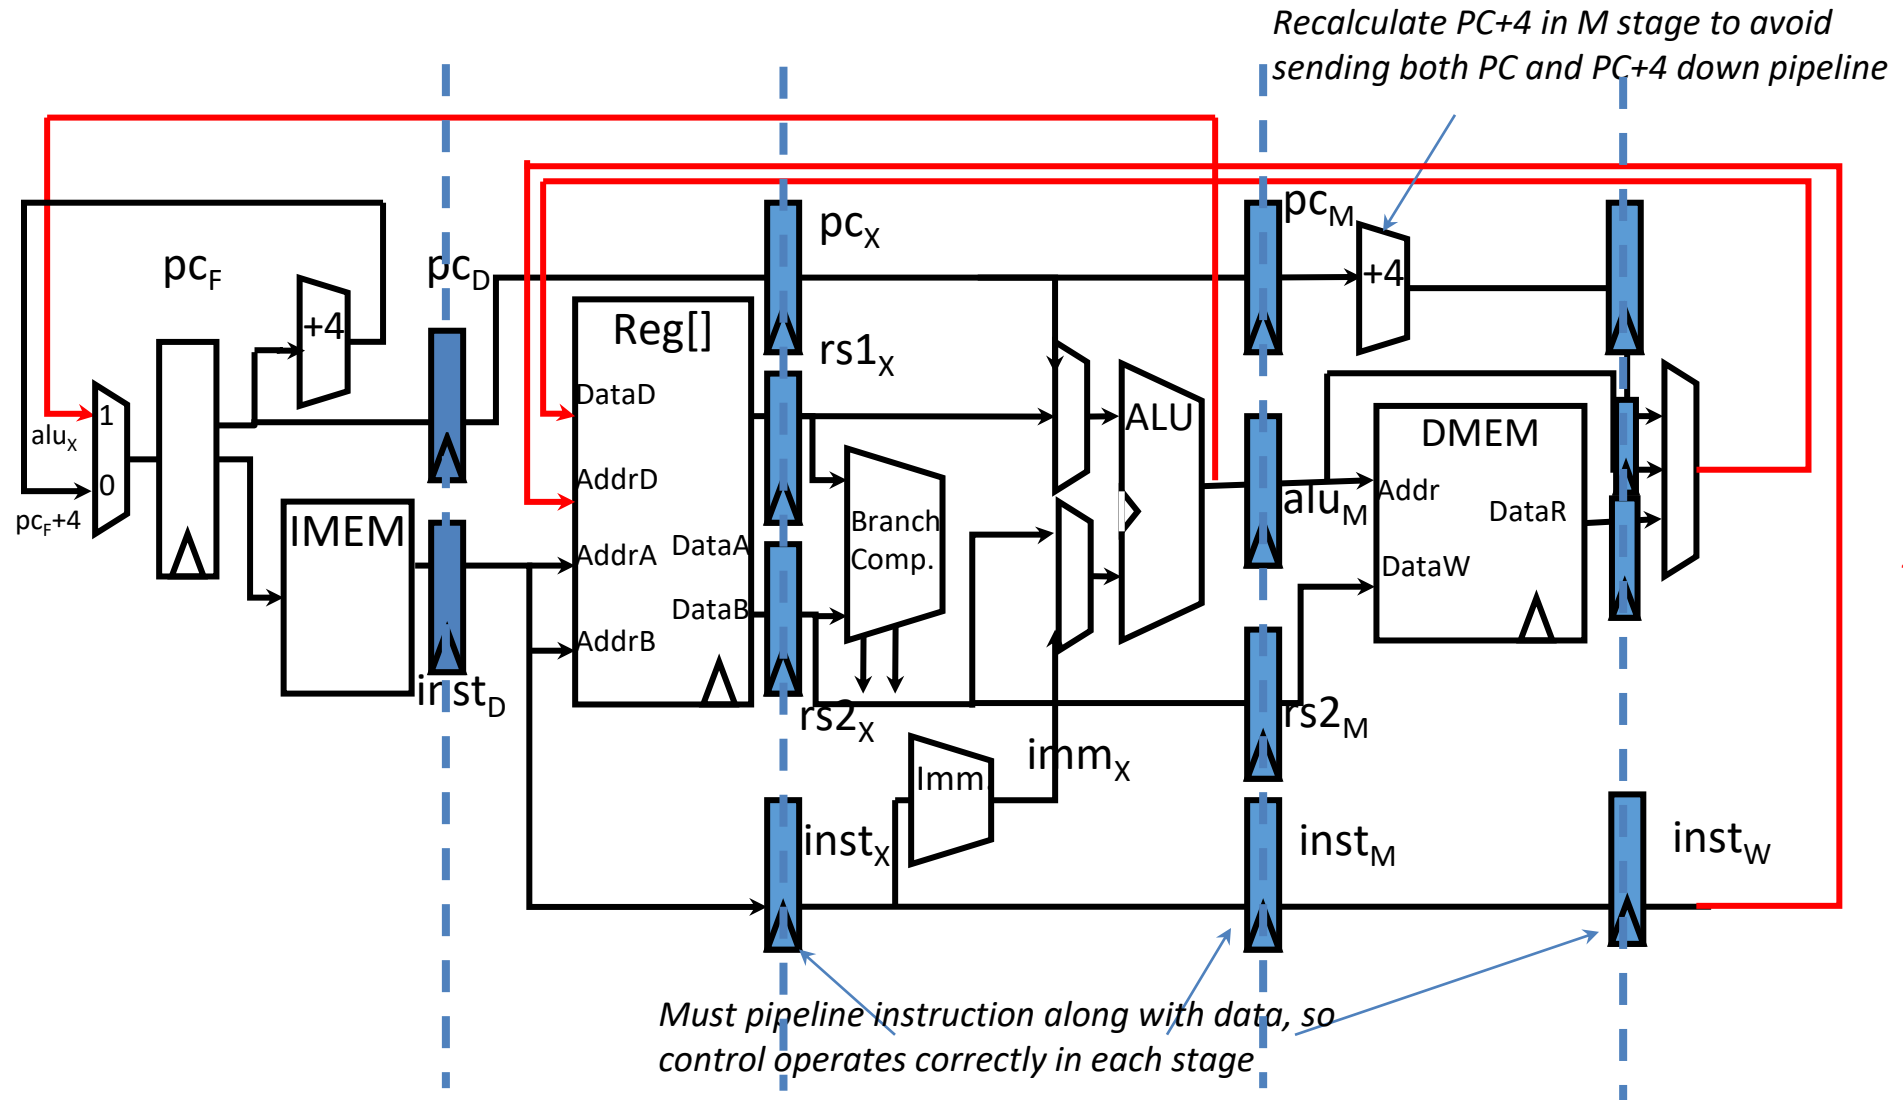- Superscalar processors

Single-Cycle RISC-V RV32I Datapath
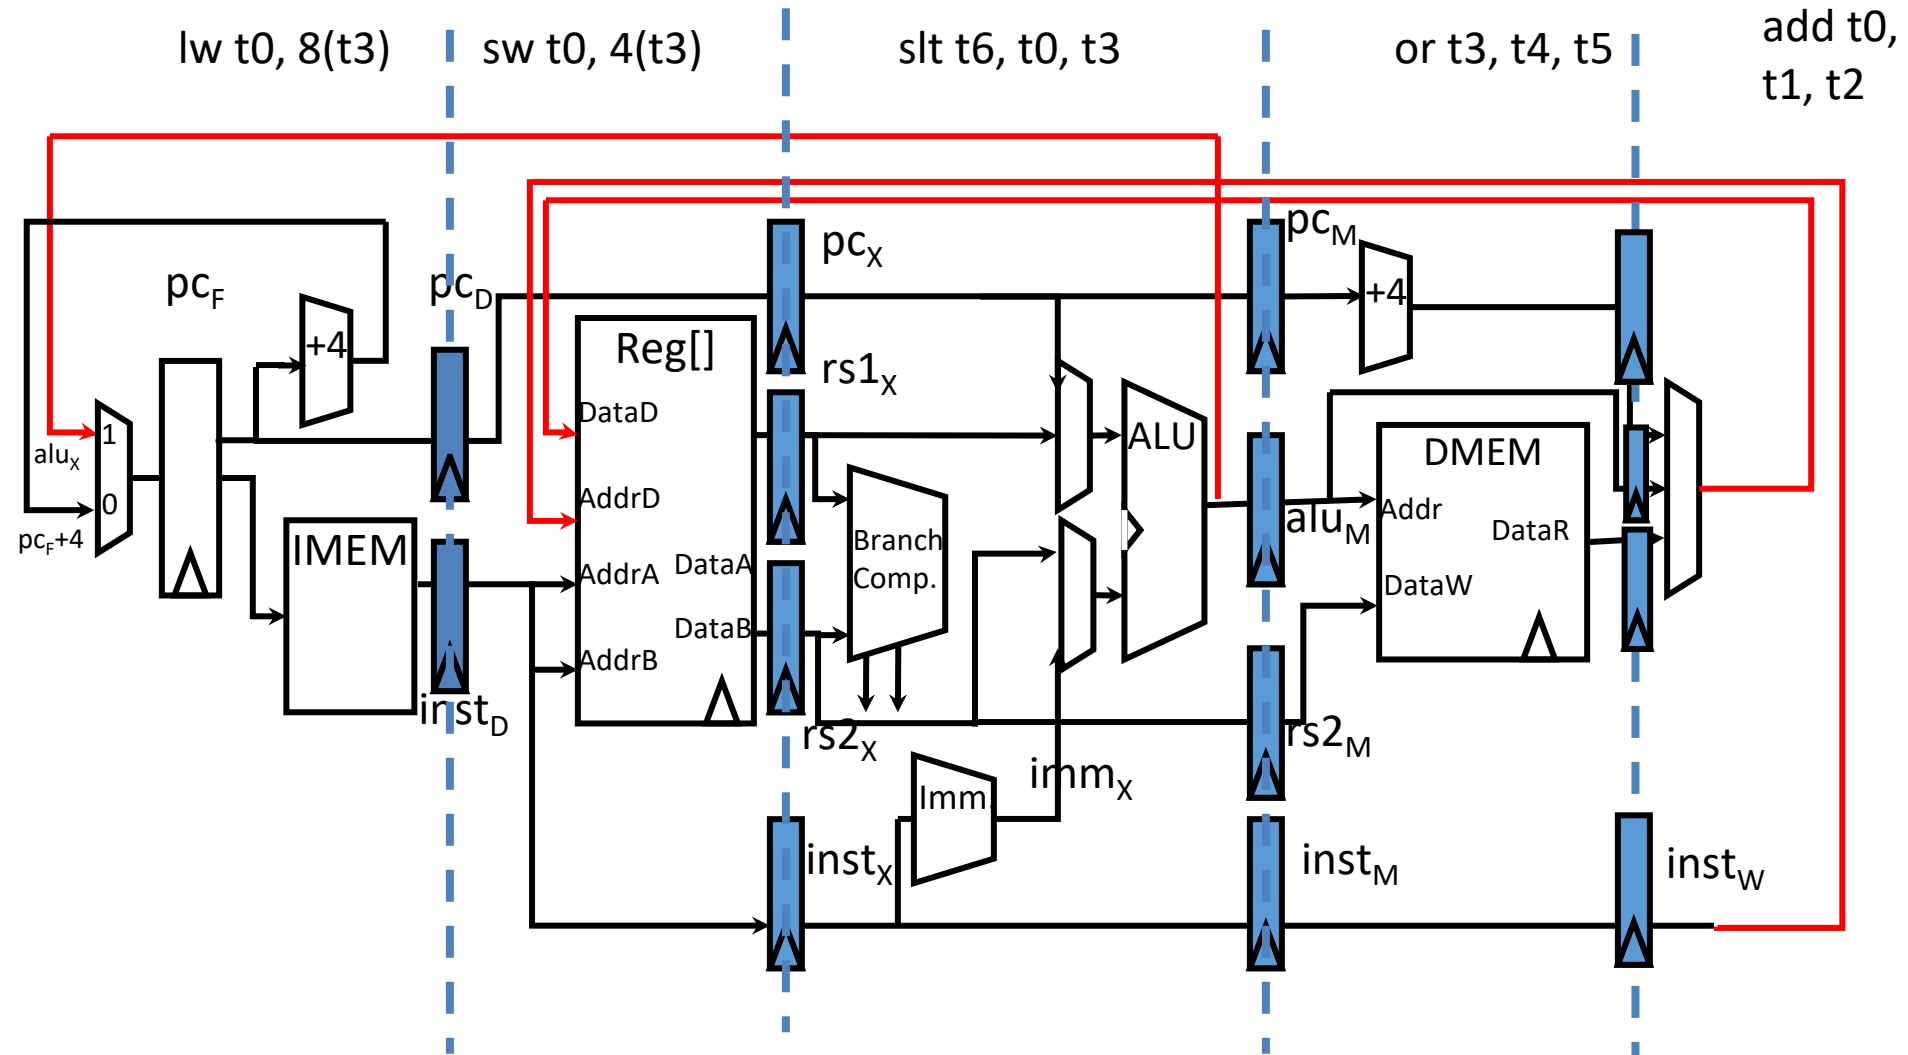
Pipelining RISC-V RV32I Datapath

Instruction Fetch (F)

Instruction Decode/Register Read (D)

ALU Execute (X)

Memory Access (M)

Write Back (W)

NOTE: Control signals are also pipelined!

Pipelined RISC-V RV32I Datapath

Recalculate PC+4 in M stage to avoid sending both PC and PC+4 down pipeline

Data from "future" stages that belong in "earlier" stages can just be wired back!

Must pipeline instruction along with data, so control operates correctly in each stage

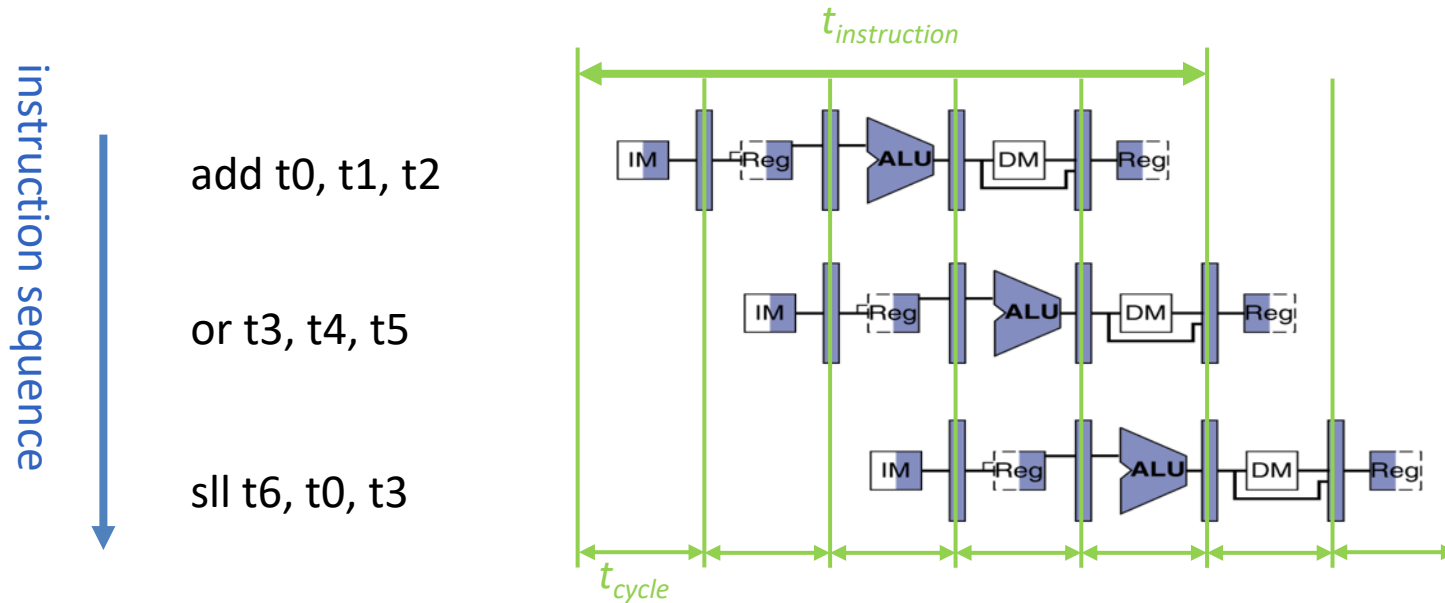# Each stage operates on different instruction

CS 61C Su20 - Lecture 14

Pipelined Control

- Control signals derived from instruction
  - As in single-cycle implementation
  - Proper Information (e.g. signals, rd) is stored in pipeline registers for use by later stages
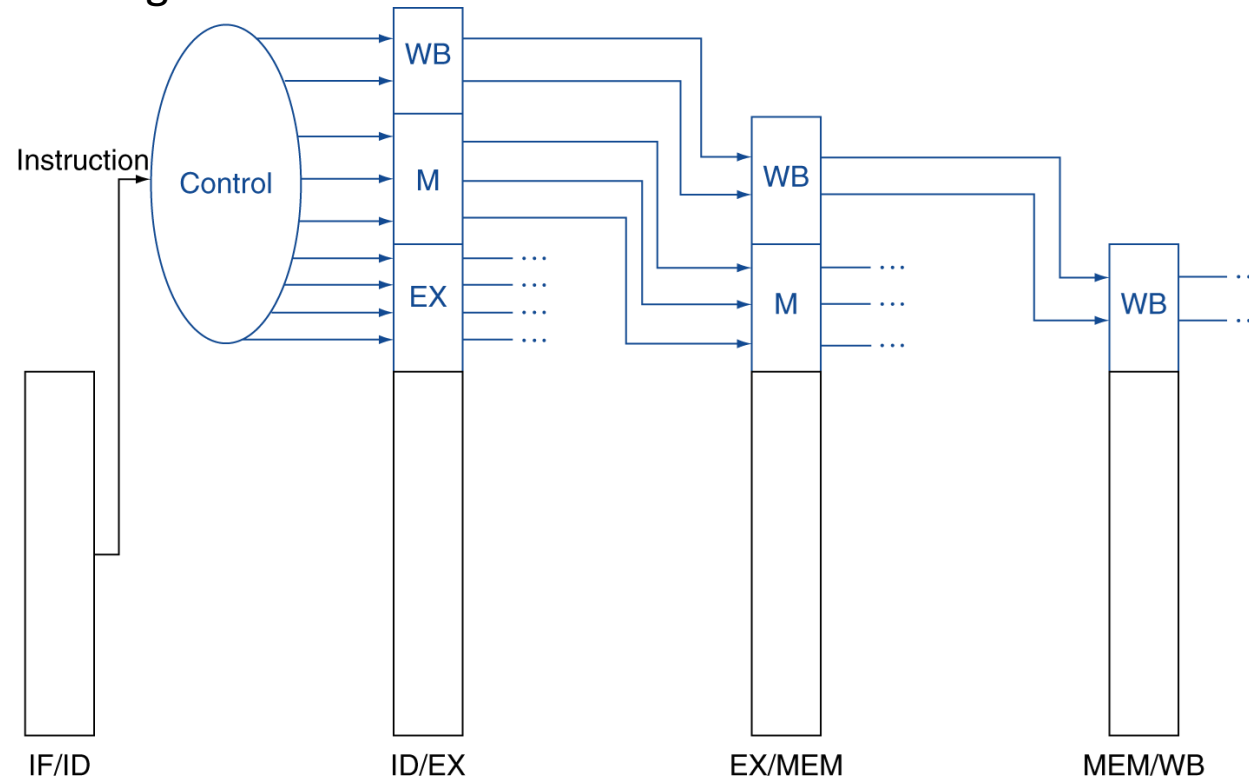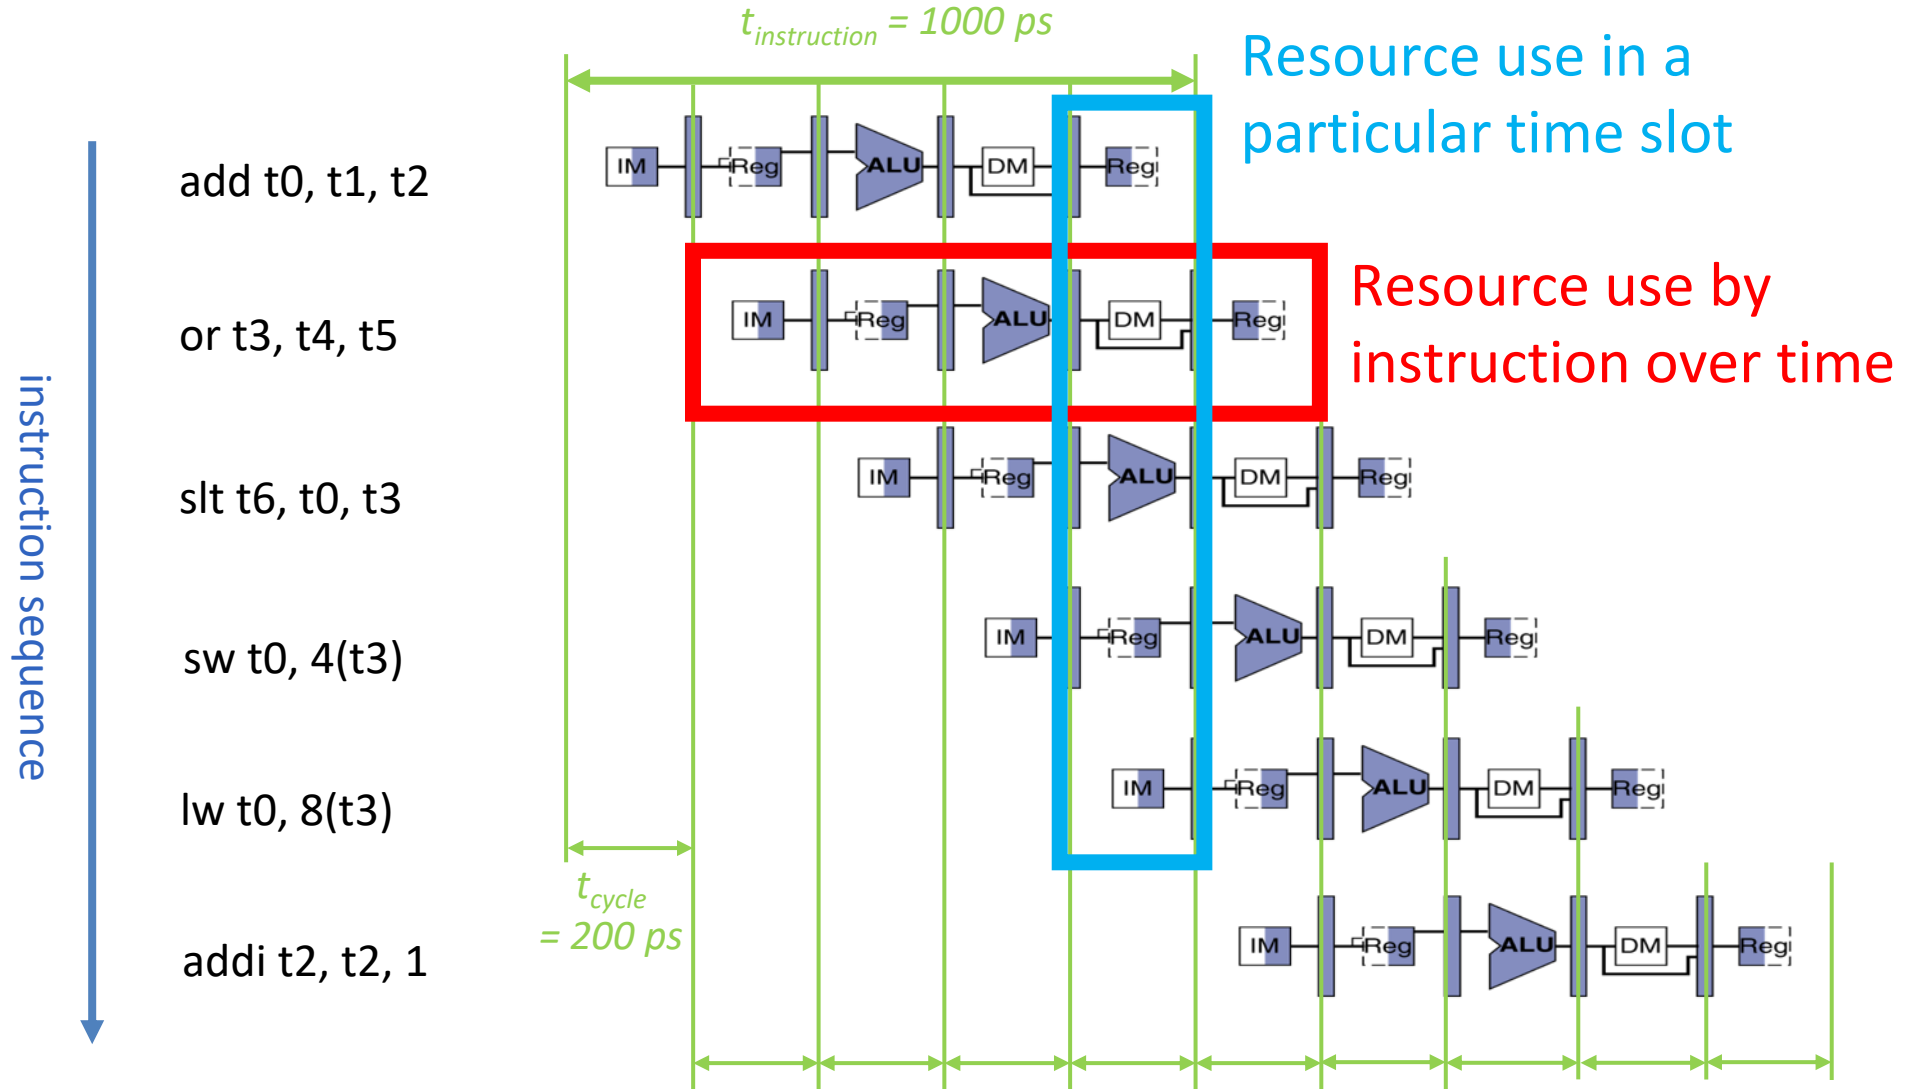
Recap: Pipelining with RISC-V

$t_{instruction}$

instruction sequence

add t0, t1, t2

or t3, t4, t5

sll t6, t0, t3

t_{cycle}

	Single Cycle	**Pipelining**
Timing	t_{step} = 100 … 200 ps	t_{cycle} = 200 ps
	Register access only 100 ps	All cycles same length
Instruction time, $t_{instruction}$	= t_{cycle} = 800 ps	1000 ps
Clock rate, f_s	1/800 ps = 1.25 GHz	1/200 ps = 5 GHz
Relative speed	1 x	4 x

RISC-V Pipeline

$t_{instruction}$ = 1000 ps

Resource use in a particular time slot

Resource use by instruction over time

instruction sequence

add t0, t1, t2

or t3, t4, t5

slt t6, t0, t3

sw t0, 4(t3)

lw t0, 8(t3)

addi t2, t2, 1

t_{cycle} = 200 ps

Question: Which of the following signals for RISC-V does NOT need to be passed into the EX pipeline stage for a beq instruction?

`beq t0 t1 Label`
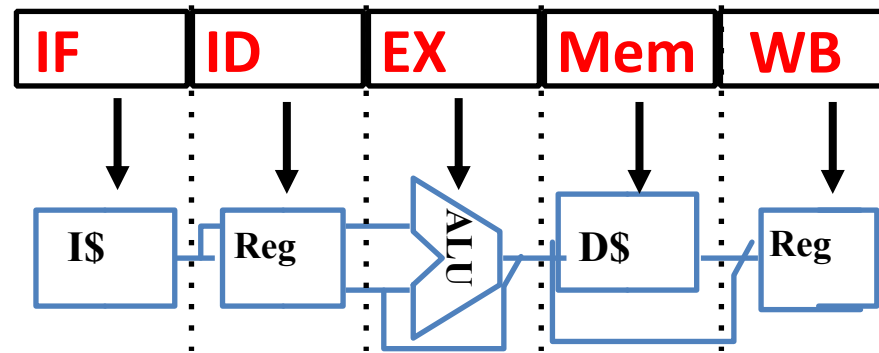
(A) **BrUn**

(B) **MemWr**

(C) **RegWr**

(D) **WBSel**

Agenda

Hazards Ahead!

- RISC-V Pipeline
- **Hazards**
  - Structural
  - Data
    - R-type instructions
    - Load
  - Control
- Superscalar processors

Pipelining Hazards

A *hazard* is a situation that prevents starting the next instruction in the next clock cycle

1) *Structural hazard*
   – A required resource is busy
     (e.g. needed in multiple stages)

2) *Data hazard*
   – Data dependency between instructions
   – Need to wait for previous instruction to complete its data write

3) *Control hazard*
   – Flow of execution depends on previous instruction
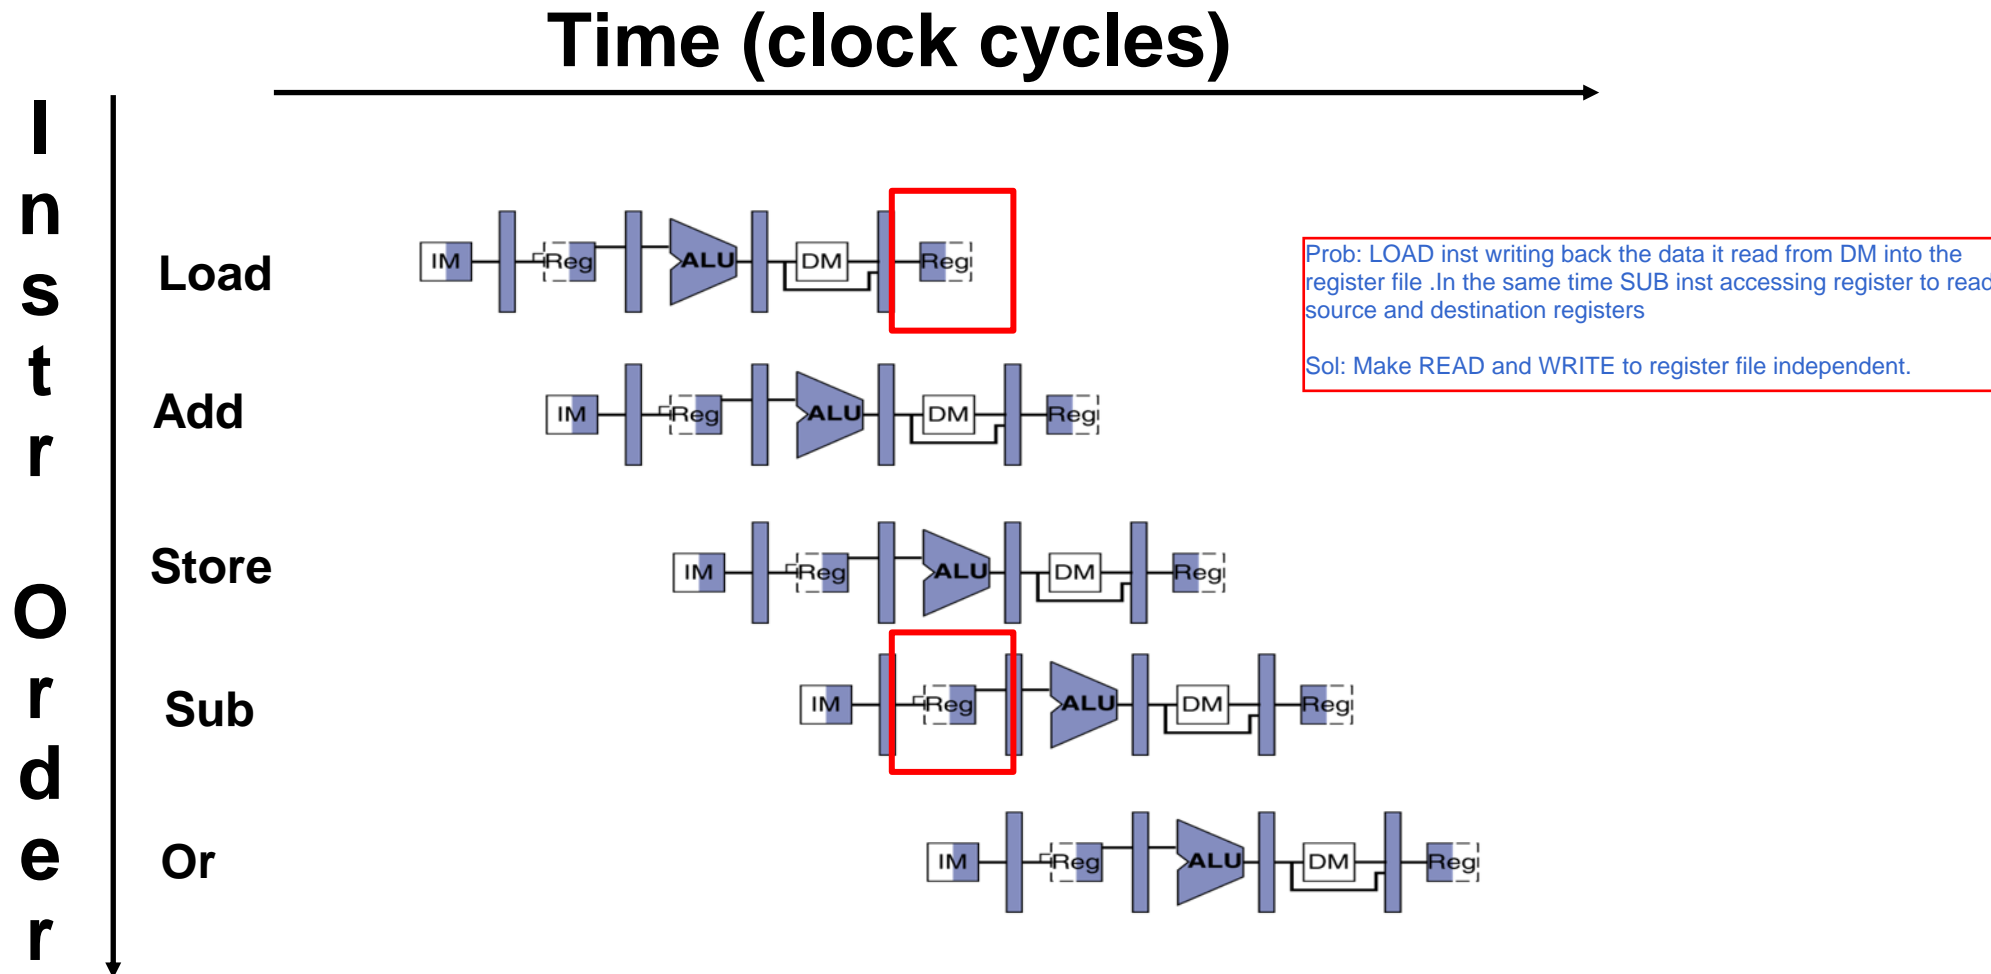
Agenda

- RISC-V Pipeline
- Hazards
  - Structural
  - Data
    - R-type instructions
    - Load
  - Control
- Superscalar processors

Structural Hazard

- **Problem:** Two or more instructions in the pipeline compete for access to a single physical resource

- **Solution 1:** Instructions take turns using resource, some instructions have to stall (wait)

- **Solution 2:** Add more hardware to machine

- Can always solve a structural hazard by adding more hardware

Structural Hazard: Regfile!

- RegFile: Used in ID and WB!

Time (clock cycles)

Prob: LOAD inst writing back the data it read from DM into the register file .In the same time SUB inst accessing register to read source and destination registers

Sol: Make READ and WRITE to register file independent.

Regfile Structural Hazards

- Each instruction:
  - can read up to two operands in decode stage
  - can write one value in writeback stage

- Avoid structural hazard by having separate "ports"
  - two independent read ports and one independent write port

- Three accesses per cycle can happen simultaneously

RW RA RB

RegWEn

5   5   5

portW

32

Clk

32 x 32-bit
Registers

portA

32

portB

32
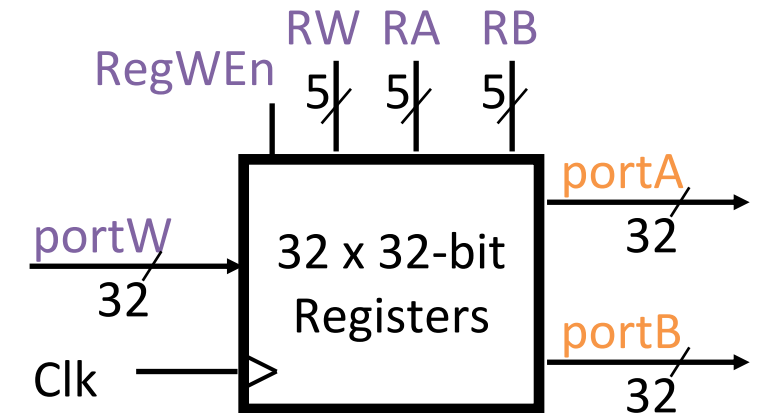
Regfile Structural Hazards

- Two *alternate* solutions:

  1) Build RegFile with independent read and write ports (what you will do in the project; good for single-stage)

  2) Double Pumping: split RegFile access in two! Prepare to write during 1st half, write on _falling_ edge, read during 2nd half of each clock cycle

     Write in negedge clock cycle and read in posedge clock cycle. Since Register file is fast we can do this

     - Will save us a cycle later...

     - Possible because RegFile access is *VERY* fast (takes less than half the time of ALU stage)

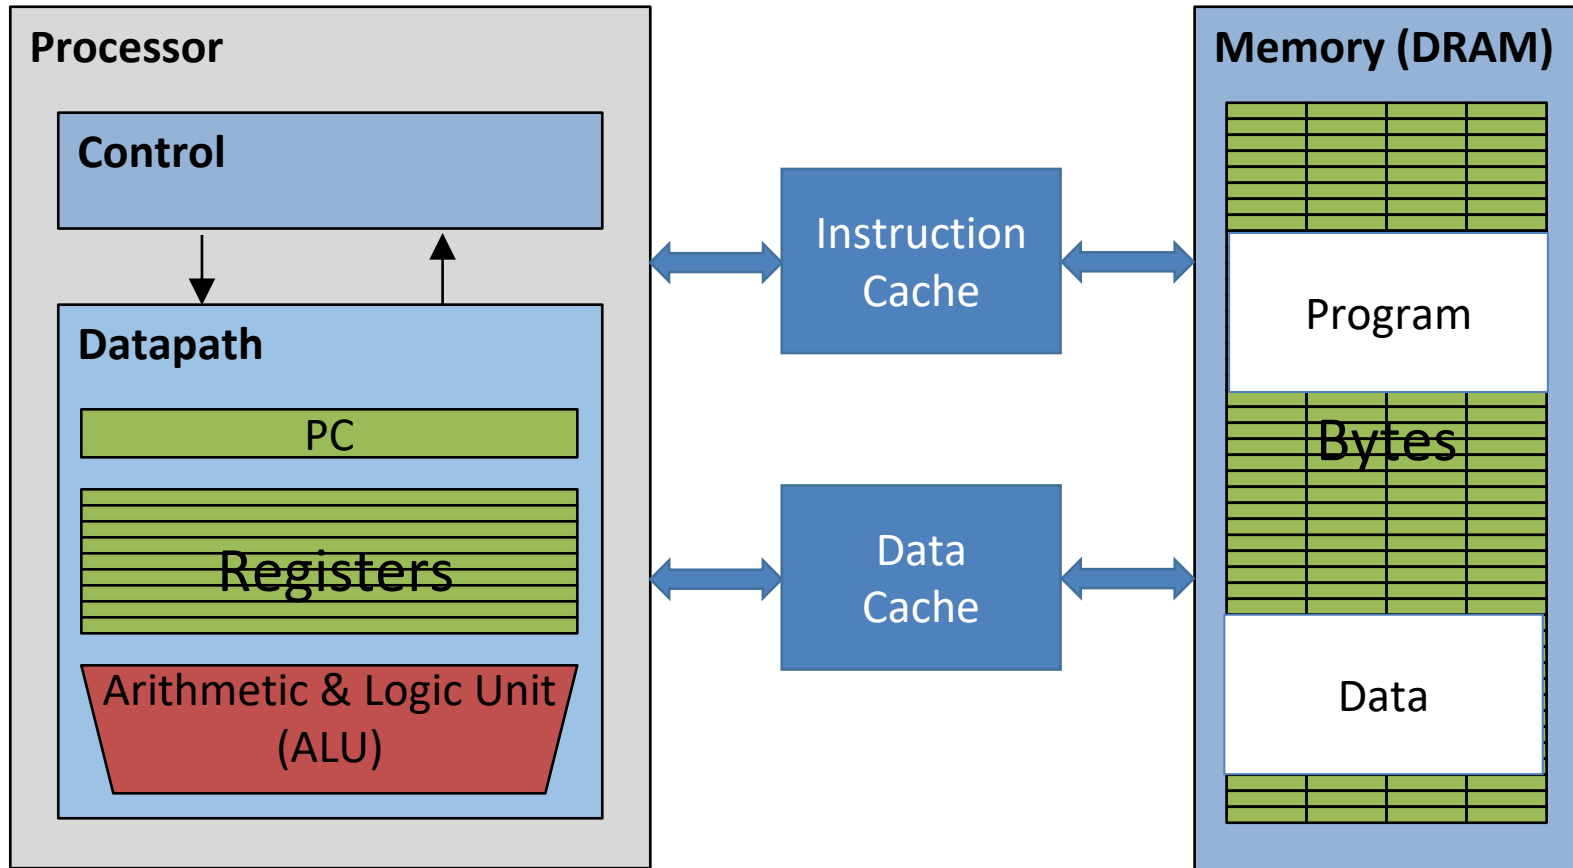- **Conclusion:** Read and Write to registers during same clock cycle is okay

Structural Hazard: Memory!

- Memory units: Used in IF and MEM!

Time (clock cycles)

Instr Order

Load

Add

Store

Sub

Or

Prob: If we use same memory to both instructions and data for the cases like this we will get conflcts.

Sol: use seperate caches for instruction and data

Instruction and Data Caches

Caches: small and fast "buffer" memories

Structural Hazards – Summary

- Conflict for use of a resource

- In RISC-V pipeline with a single memory unit
  - Load/store requires data access
  - Without separate memory units, instruction fetch would have to *stall* for that cycle
    - All other operations in pipeline would have to wait

- Pipelined datapaths require separate instruction/data memory units
  - Or separate instruction/data caches

- RISC ISAs (including RISC-V) designed to avoid structural hazards
  - e.g. at most one memory access/instruction

Agenda

- RISC-V Pipeline
- Hazards
  - Structural
  - Data
    - R-type instructions
    - Load
  - Control
- Superscalar processors

2. Data Hazards (1/2)

- Consider the following sequence of instructions:

```
add   t0, t1, t2
sub   t4, t0, t3
and   t5, t0, t6
or    t7, t0, t8
xor   t9, t0, t10
```
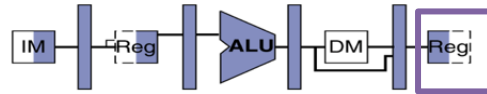
Stored during WB

Read during ID
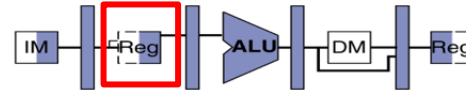
2. Data Hazards (2/2)

Identifying data hazards:

- Where is data **WRITTEN**?

- Where is data **READ**?

- Does the WRITE happen AFTER the READ?

Time (clock cycles)
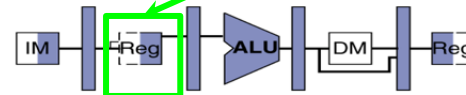
add **t0**, t1, t2

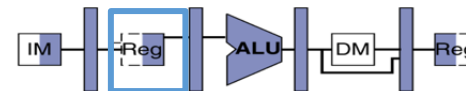sub t4, **t0**, t3

and t5, **t0**, t6

or   t7, **t0**, t8
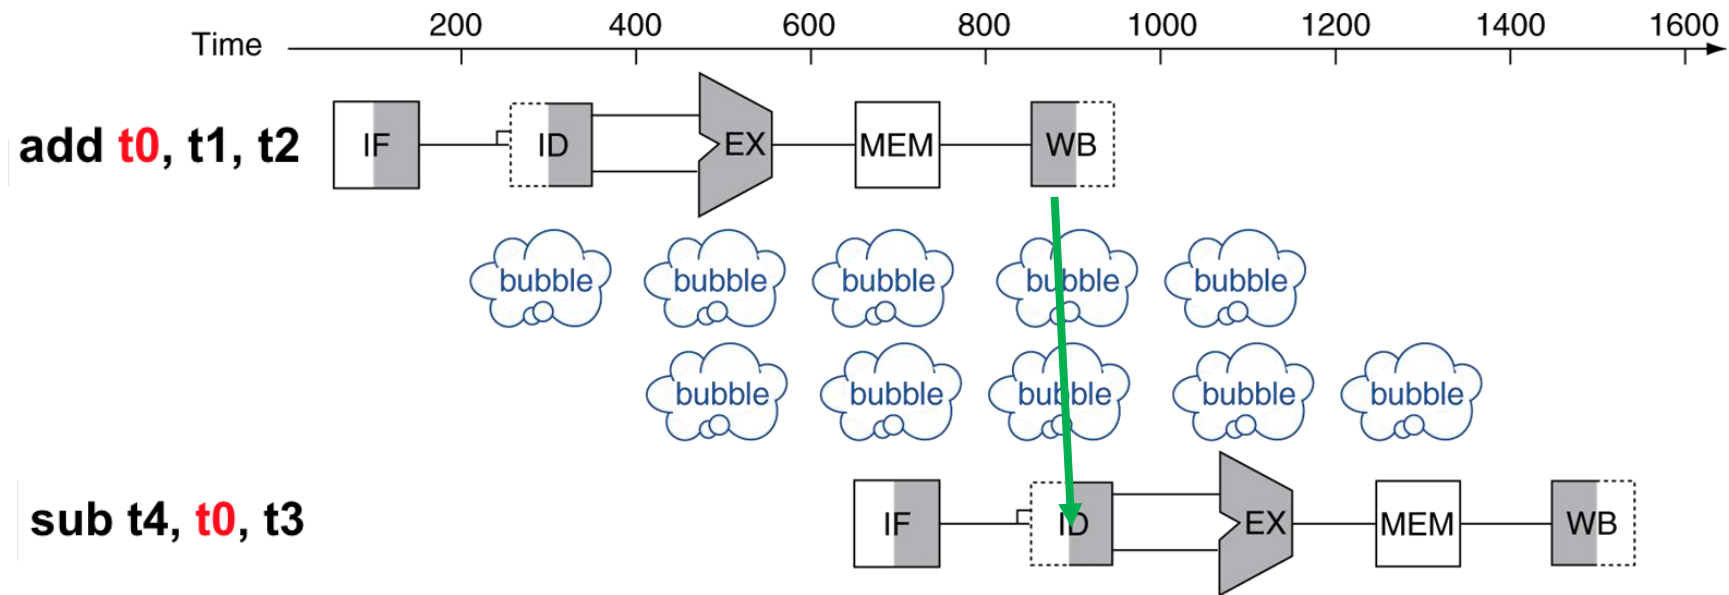
xor t9, **t0**, t10

Only an issue if no double pumping!

Solution 1: Stalling

- Problem: Instruction depends on result from previous instruction
    - add   t0, t1, t2
      sub   t4, t0, t3

- Bubble:
    - effectively NOP: affected pipeline stages do "nothing" (add x0 x0 x0)

Stalls and Performance

- Stalls reduce performance
  - Decrease throughput of "valid" or useful instructions
  - Can also be seen as increasing the latency of our stalled instruction
- But stalls are required to get correct results
- Compiler can arrange code to avoid hazards and stalls!
  - And so can 61C students ;)
  - Requires knowledge of the pipeline structure, and knowledge of instruction interactions

Data Hazard Solution: Forwarding

- Forward result as soon as it is available, even though it's not stored in RegFile yet
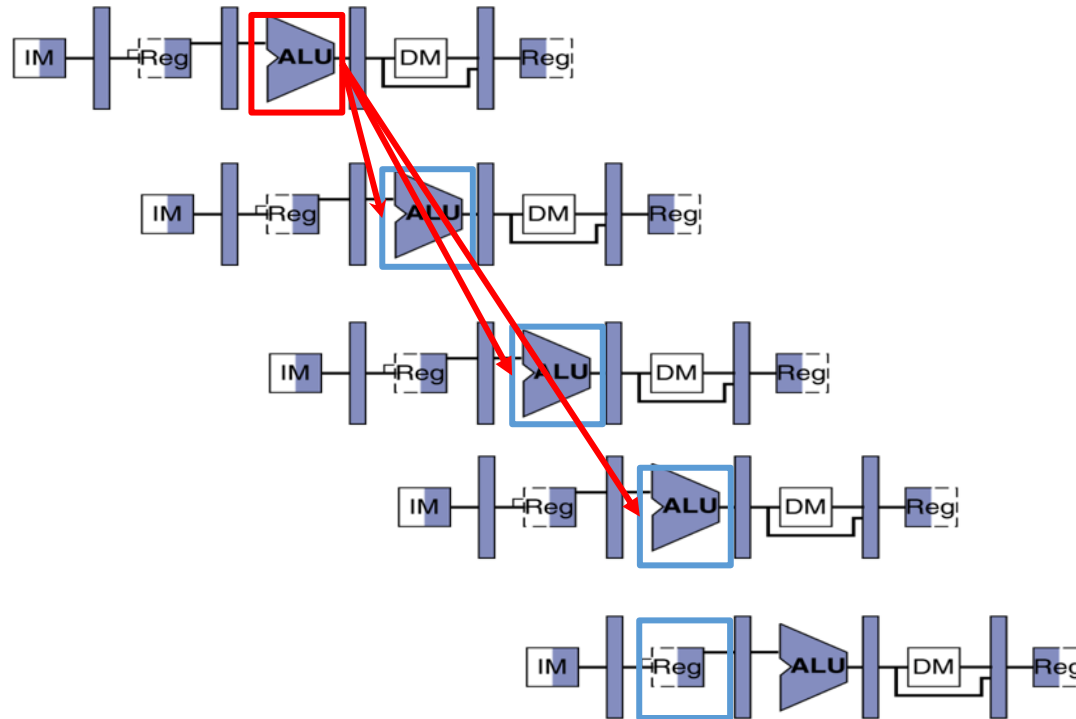
add   t0, t1, t2
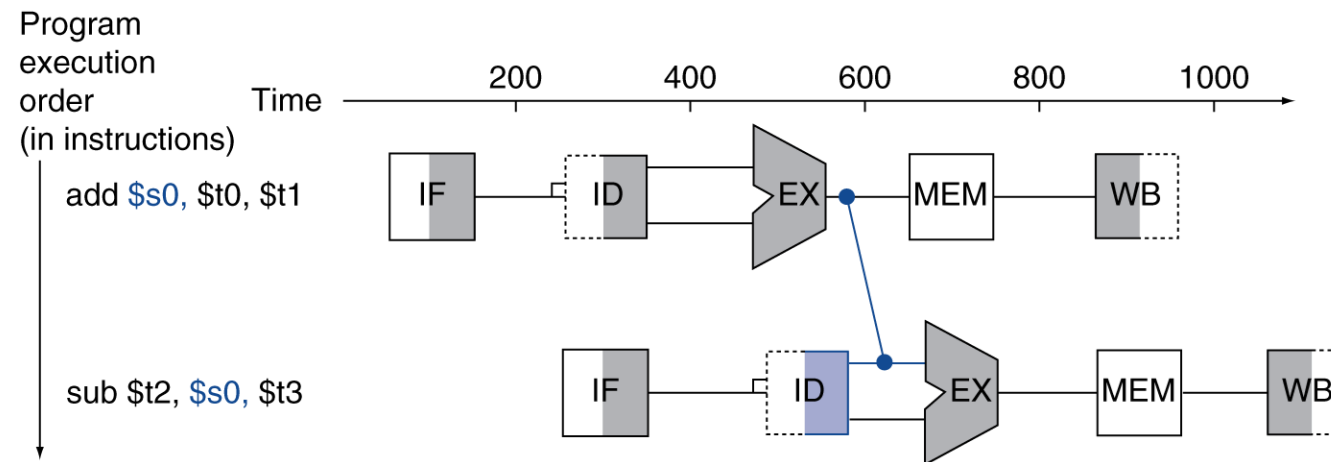
sub   t4, t0, t3

and   t5, t0, t6

or   t7, t0, t8

xor   t9, t0, t10

Forwarding: grab operand from pipeline stage, rather than register file

Forwarding (aka Bypassing)

- Use result when it is computed
  - Don't wait for it to be stored in a register
  - Requires extra hardware in the datapath (and extra control!)
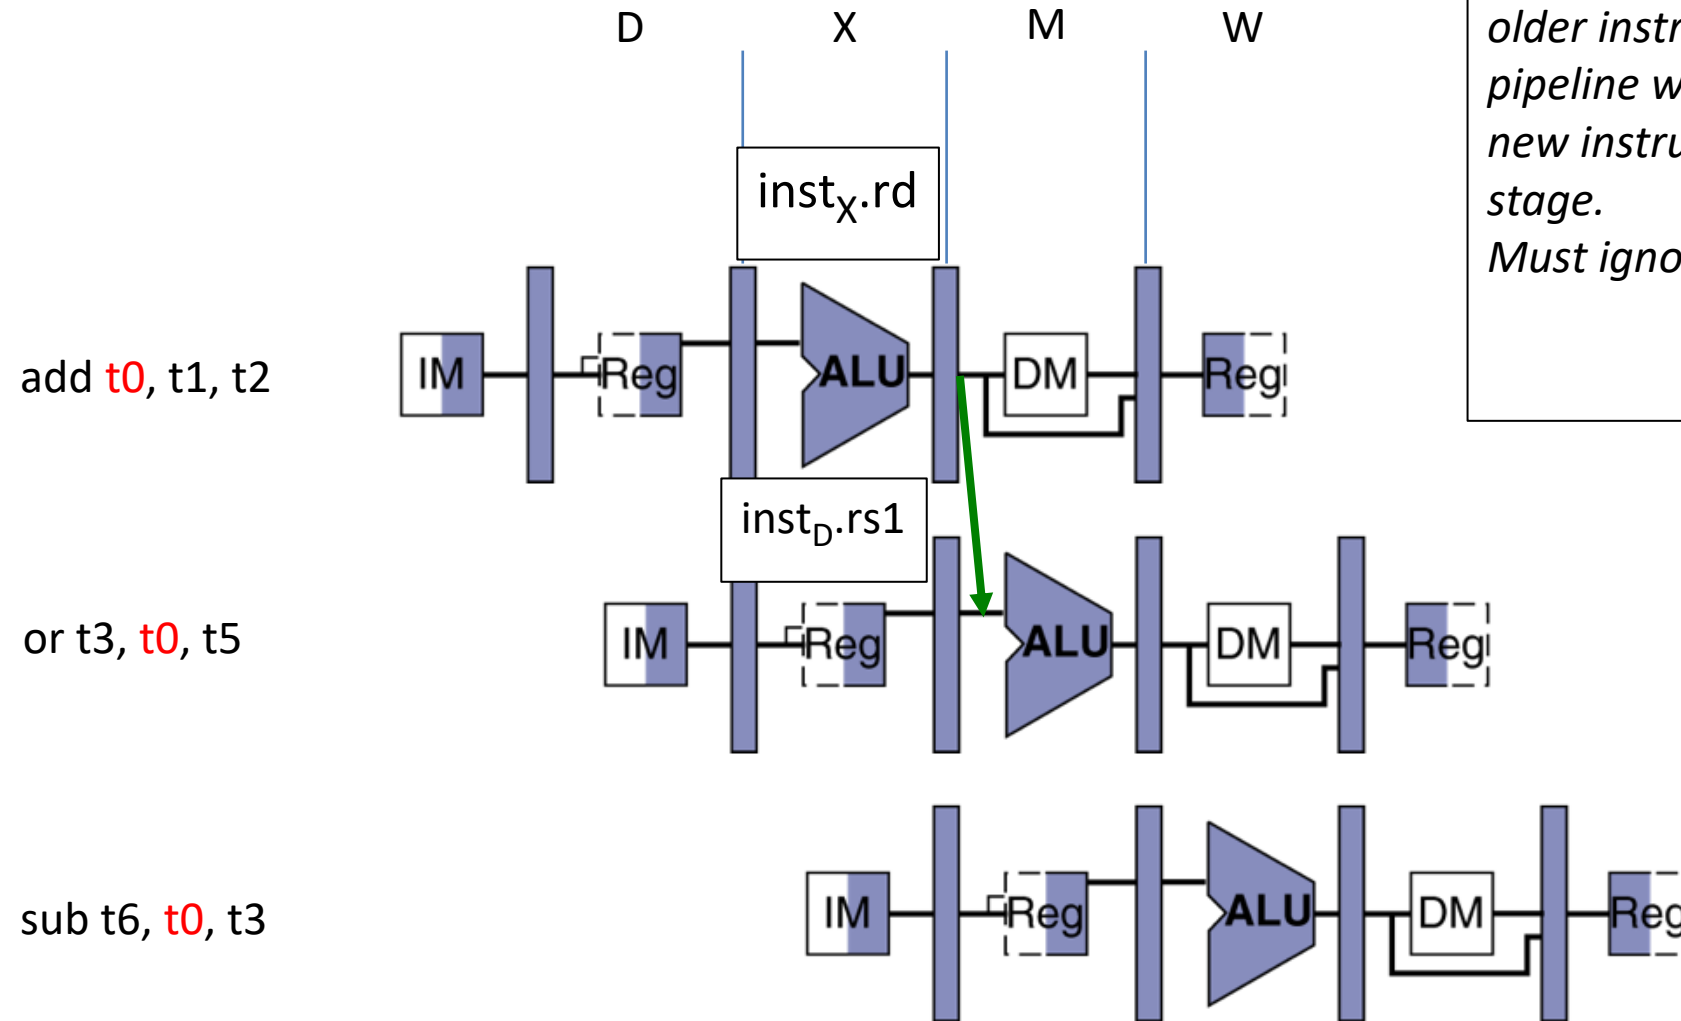 - Not required on project 3 :)

Question

In our 5-stage pipeline, how many subsequent instructions do we need to look at to detect data hazards for this **add**? Assume we have double-pumping.

A) 1 instruction
B) 2 instructions
C) 3 instructions
D) 4 instructions
E) 5 instructions

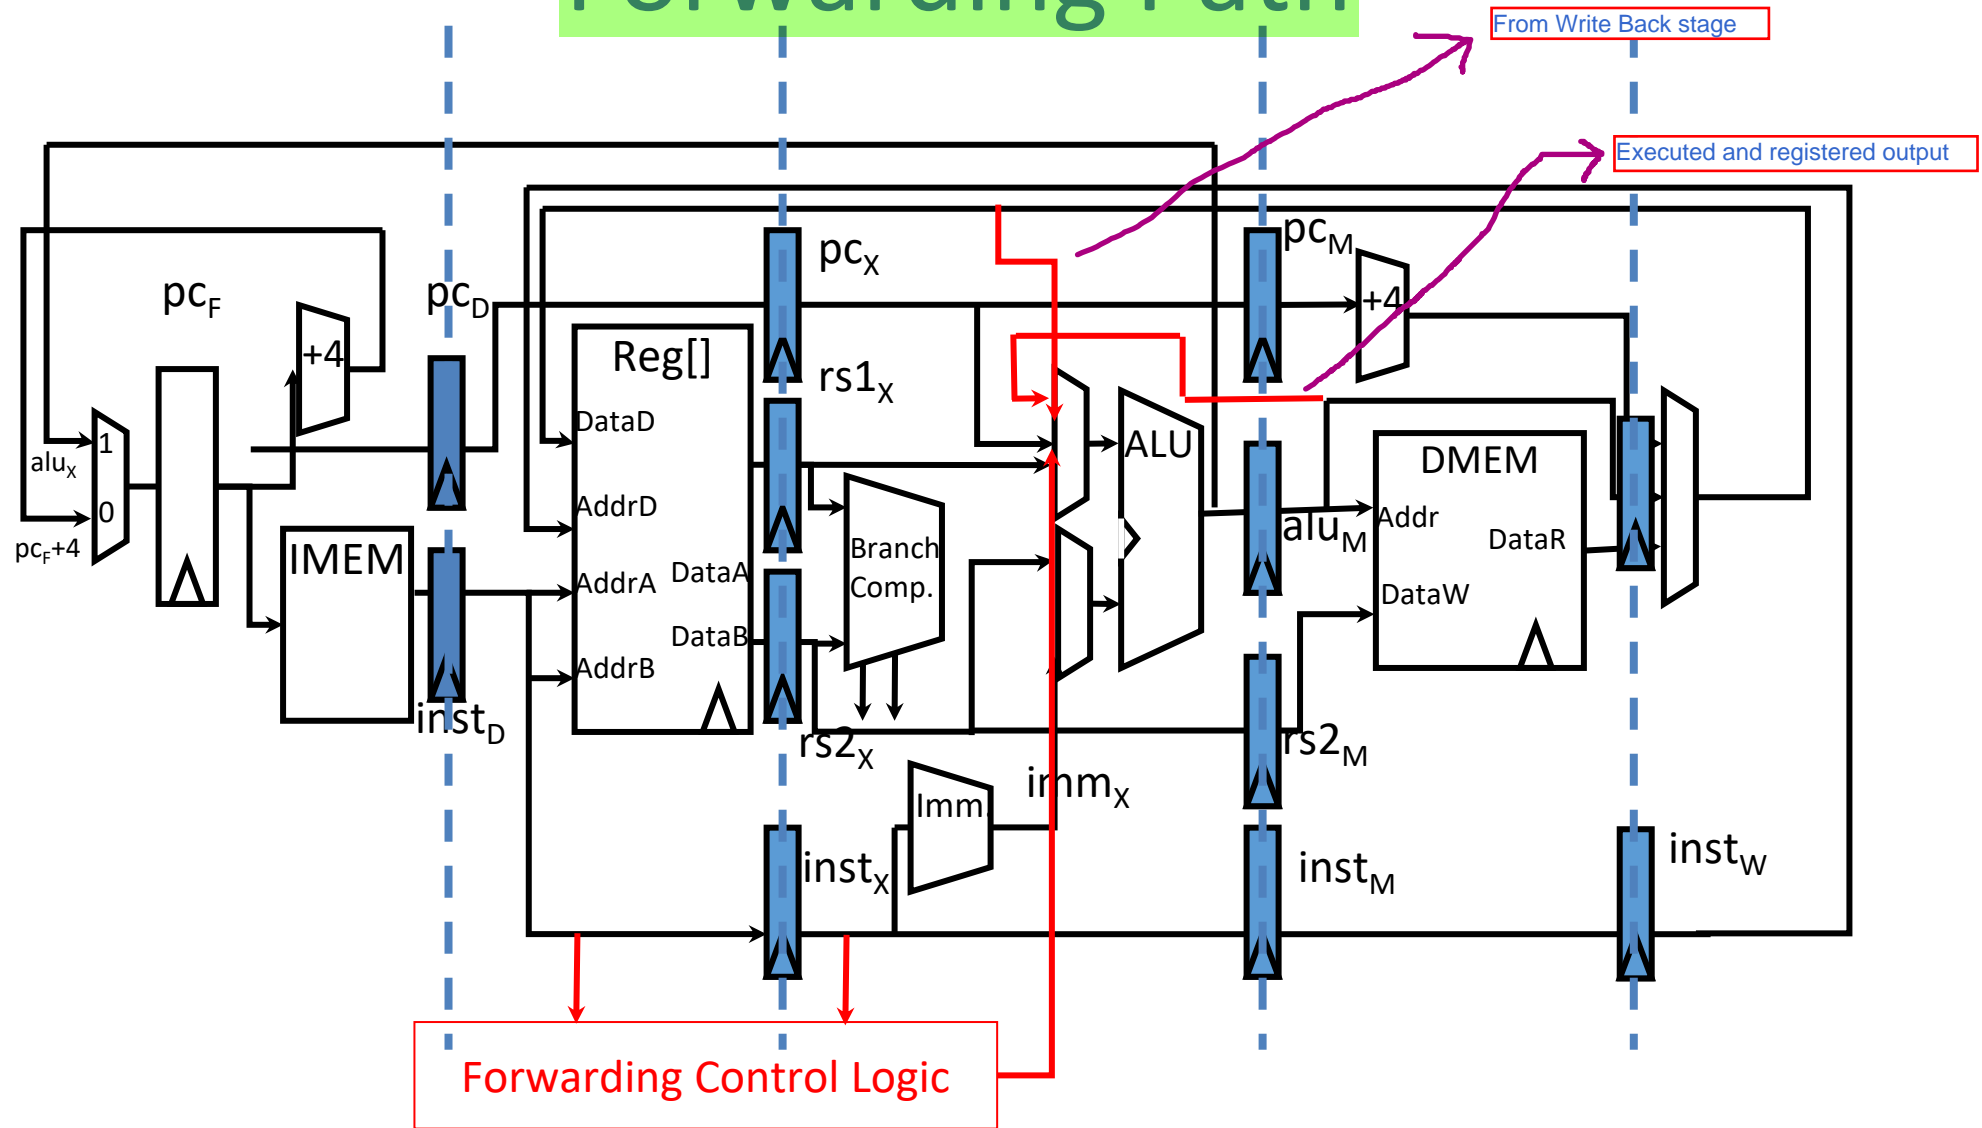add **t0**, t1, t2

sub t4, **t0**, t3

and t5, **t0**, t6

or   t7, **t0**, t8

xor t9, **t0**, t10
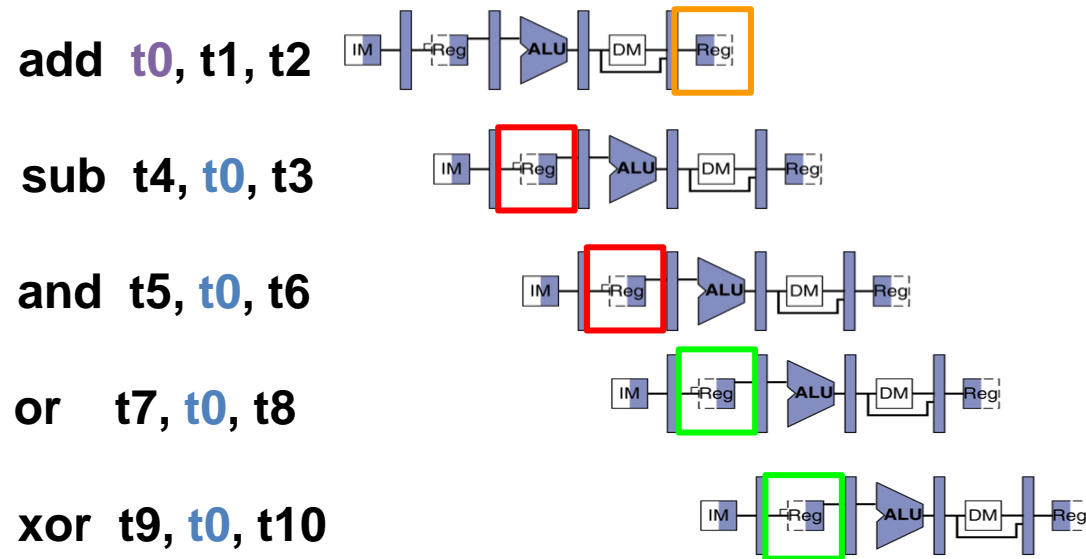
Detect Need for Forwarding (example)

D   X   M   W

$inst_X.rd$

$inst_D.rs1$

add t0, t1, t2

or t3, t0, t5

sub t6, t0, t3

Compare destination of older instructions in pipeline with sources of new instruction in decode stage.
Must ignore writes to x0!

Forwarding Path

Agenda

- RISC-V Pipeline
- Hazards
  - Structural
  - Data
    - R-type instructions
    - Load
  - Control
- Superscalar processors

Data Hazard: Loads (1/4)

- **Recall:** Dataflow backwards in time are hazards

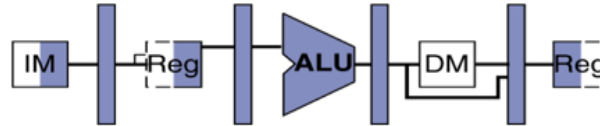**lw t0, 0(t1)**

sub t3, t0, t2

- Can't solve all cases with forwarding
  - Must *stall* instruction <u>dependent</u> on load (sub), then forward after the load is done (more hardware)
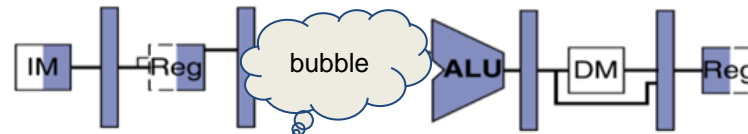
Data Hazard: Loads (2/4)

- *Hardware* stalls pipeline
  - Called "hardware interlock"

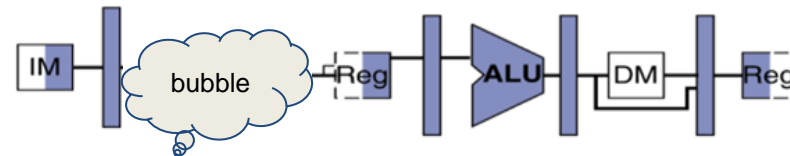This is what happens in hardware in a "hardware interlock"

lw **t0**, 0(t1)

sub t3, t0, t2

and t5, t0, t4

or   t7, t0, t6
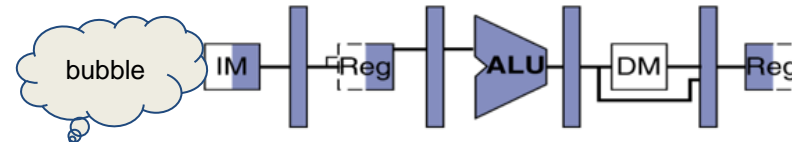
Data Hazard: Loads (3/4)

- Stall is equivalent to `nop`

lw **t0**, 0(t1)

nop

sub t3, **t0**, t2

and t5, t0, t4

or   t7, t0, t6
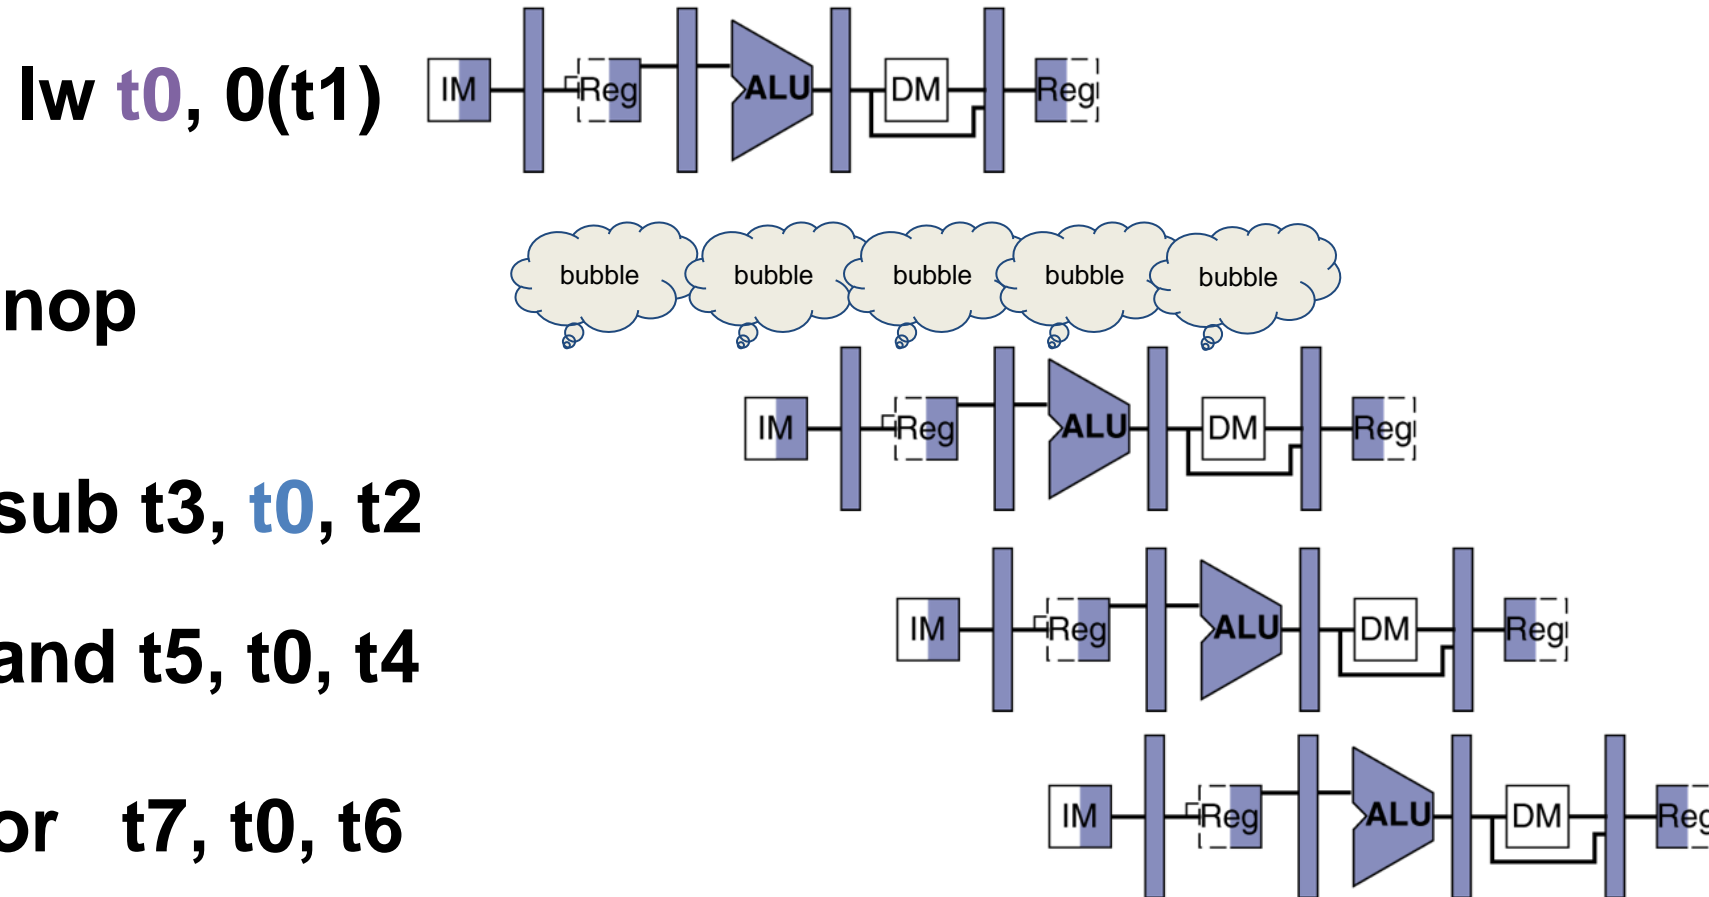
Data Hazard: Loads (4/4)

- Slot after a load is called a *load delay slot*
  - If that instruction uses the result of the load, then the hardware will stall for <u>one cycle</u>
  - Equivalent to inserting an explicit `nop` in the slot
    - except the latter uses more code space
  - Performance loss
- **Idea:** Let the compiler/assembler put an unrelated instruction in that slot → no stall!

Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction!

- RISC-V code for `D=A+B;  E=A+C;`

t3 t1   t2   t5   t4

Original Order: **13 cycles**
```
lw t1, 0(t0)
lw t2, 4(t0)
add t3, t1, t2
sw t3, 12(t0)
lw t4, 8(t0)
add t5, t1, t4
sw t5, 16(t0)
```
Stall!
Stall!

Alternative: **11 cycles**
```
lw t1, 0(t0)
lw t2, 4(t0)
lw t4, 8(t0)
add t3, t1, t2
sw t3, 12(t0)
add t5, t1, t4
sw t5, 16(t0)
```

Agenda

- RISC-V Pipeline
- Hazards
  - Structural
  - Data
    - R-type instructions
    - Load
  - Control
- Superscalar processors

3. Control Hazards

- Branch (`beq, bne, ...`) determines flow of control
  - Fetching next instruction <u>depends on branch outcome</u>
  - Pipeline can't always fetch correct instruction
    - Result isn't known until end of execute

- **Simple Solution:** Stall on *every* branch until we have the new PC value
  - How long must we stall?

Branch Stall

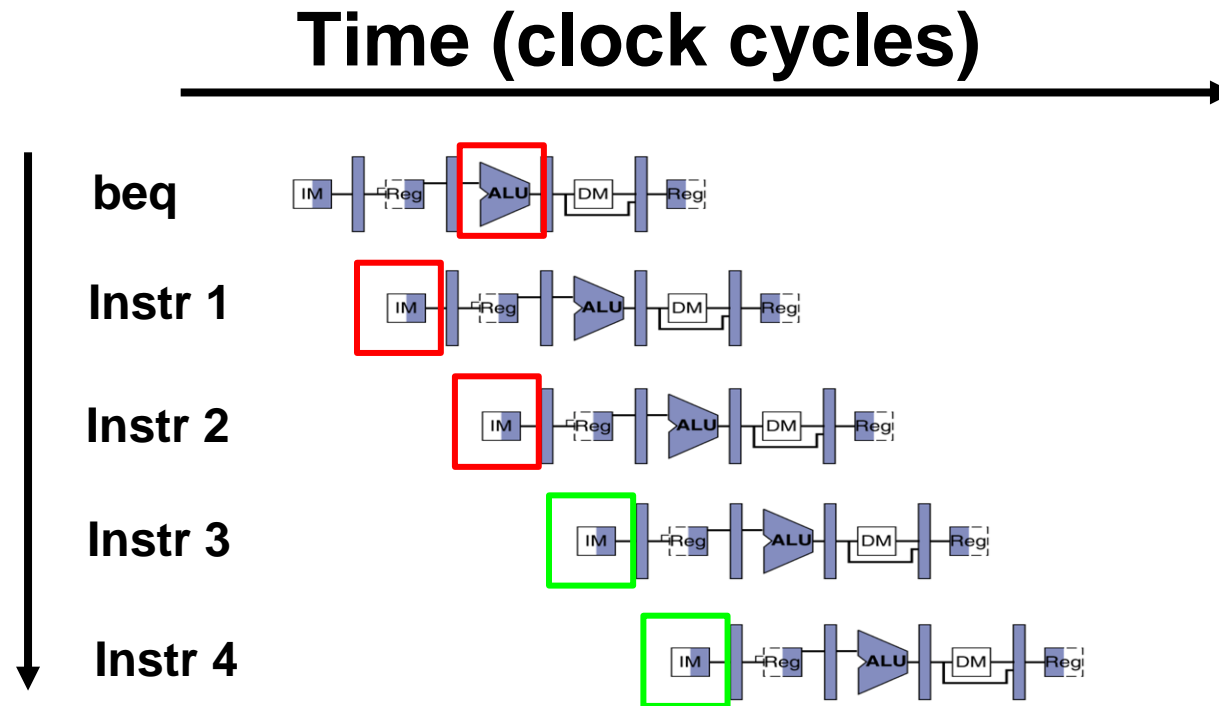- How many bubbles are required to account for the control hazard from **beq**?

A) 1

B) 2

C) 3

D) 4

E) 5

Time (clock cycles)

3. Control Hazard: Branching

- **Option #1:** Move **branch comparator** to ID stage
  - As soon as instruction is decoded, immediately make a decision and set the new value of PC
  - **Benefit:** Branch decision made in 2nd stage, so only one `nop` is needed instead of two
  - **Side Note:** Have to compute new PC value (PC + imm) in ID instead of EX
    - Adds extra copy of new-PC logic in ID stage
    - Branches are idle in EX, MEM, and WB

Improved Branch Stall

- When is comparison result available?

Time (clock cycles)

beq

Instr 1

Only one stall needed!
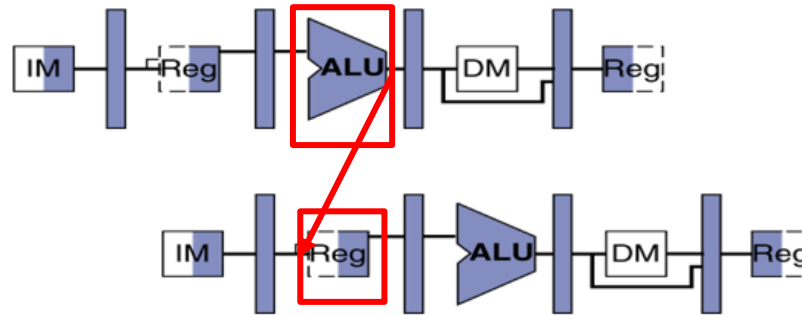
Instr 2

Instr 3

Instr 4
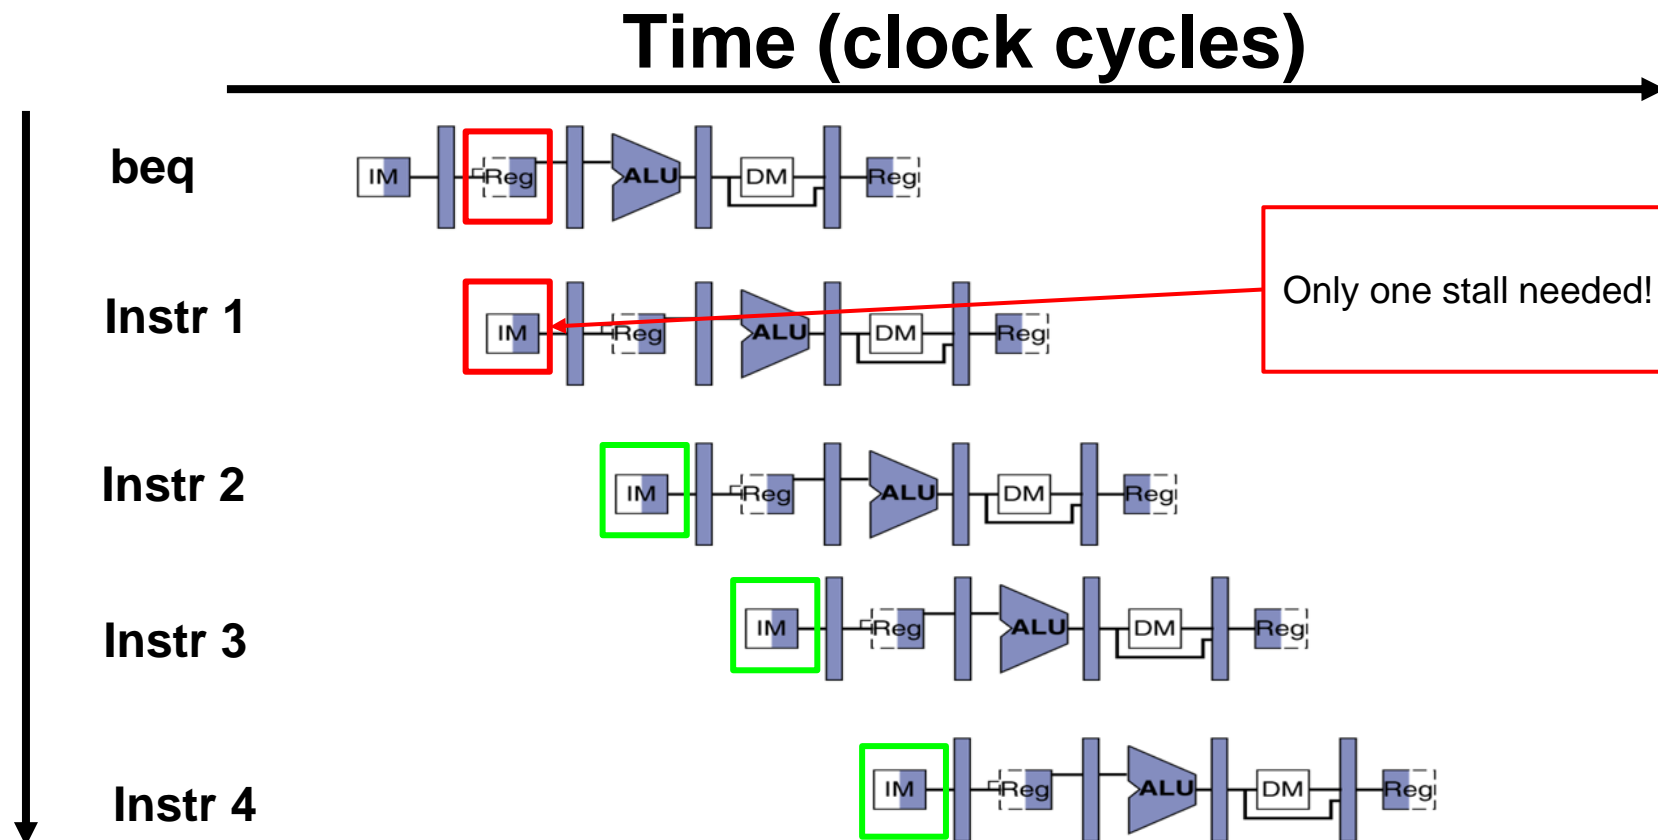
Data Hazard: Branches!

- **Recall:** Dataflow backwards in time are hazards

add t0, t0, t1

beq x0, t0, foo

- Now that t0 is needed earlier (ID instead of EX), we can't forward it to the beq's ID stage
  - Must *stall* after add, then forward (more hardware)

Observations

- **Takeaway**: Moving branch comparator to ID stage would add redundant hardware and introduce new problems

- Can we work with the nature of branches?
  - condition to jump is not satistfied
  - If branch not taken, then instructions fetched sequentially after branch are correct
  - condition to jump is satistfied
  - If branch or jump taken, then need to flush incorrect instructions from pipeline by converting to NOPs

Agenda

- Structural Hazards
- Data Hazards
  - Forwarding
- Data Hazards (Continued)
  - Load Delay Slot
- **Control Hazards**
  - Branch and Jump Delay Slots
  - **Branch Prediction**

Kill Instructions after Branch if Taken

beq t0, t1, label — Taken branch

sub t2, s0, t5 — Convert to NOP

or t6, s0, t3 — Convert to NOP

label: xxxxxx — PC updated reflecting branch outcome
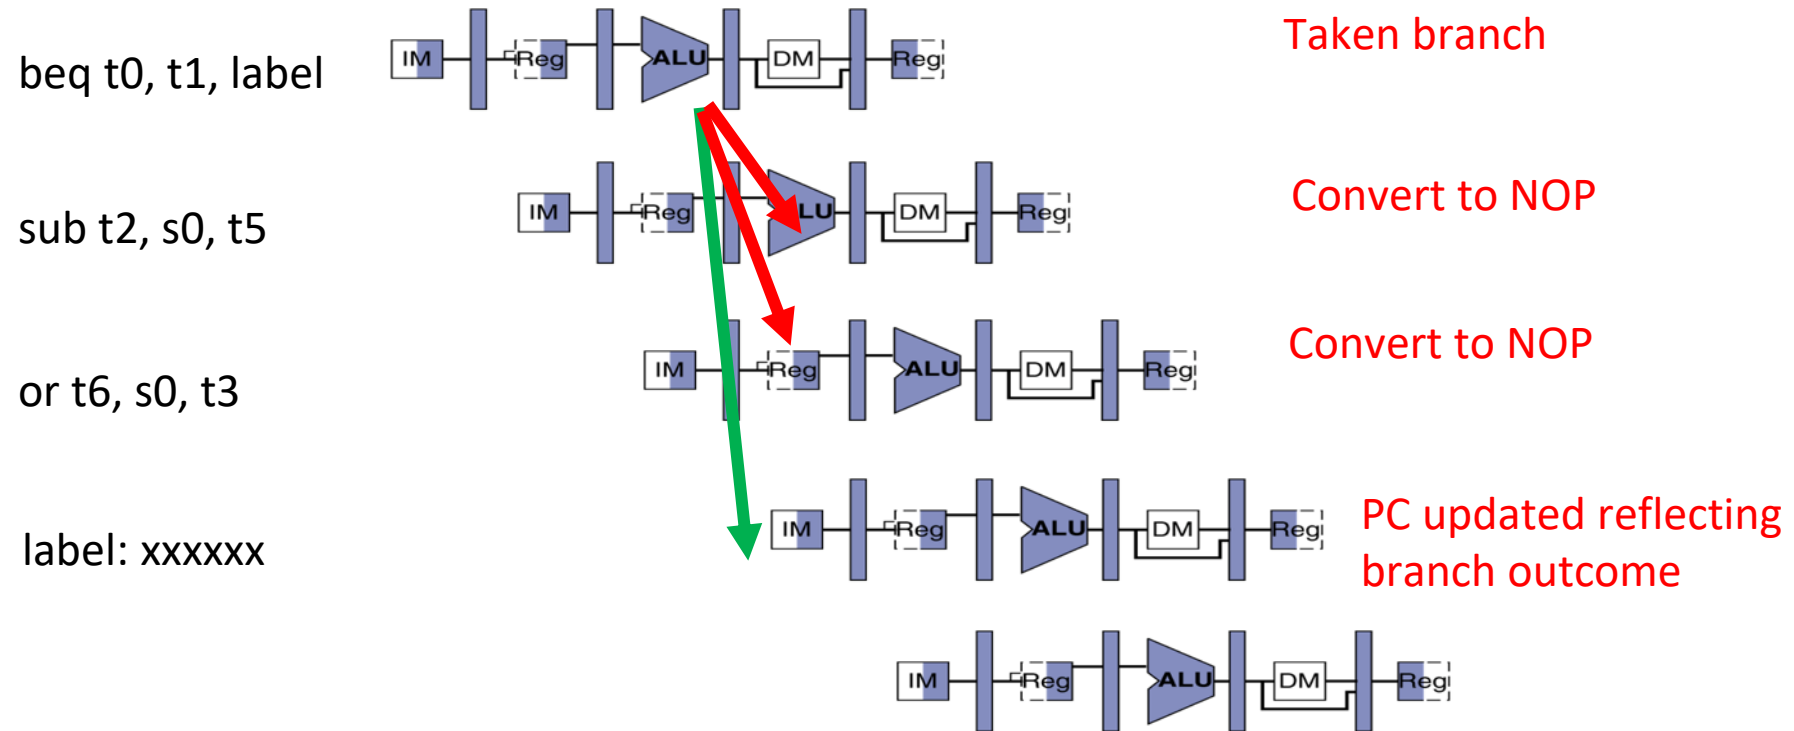
Two instructions are affected by an incorrect branch, just like we'd have to insert two NOP's/stalls in the pipeline to wait on the correct value!

3. Control Hazard: Branching

- **RISC-V Solution:** *Branch Prediction* – guess outcome of a branch, fix afterwards if necessary
  - Must cancel (*flush*) all instructions in pipeline that depended on guess that was wrong
  - How many instructions do we end up flushing?

Branch Prediction

beq t0, t1, label

label: .....

.....

Taken branch

Guess next PC!

Check guess correct
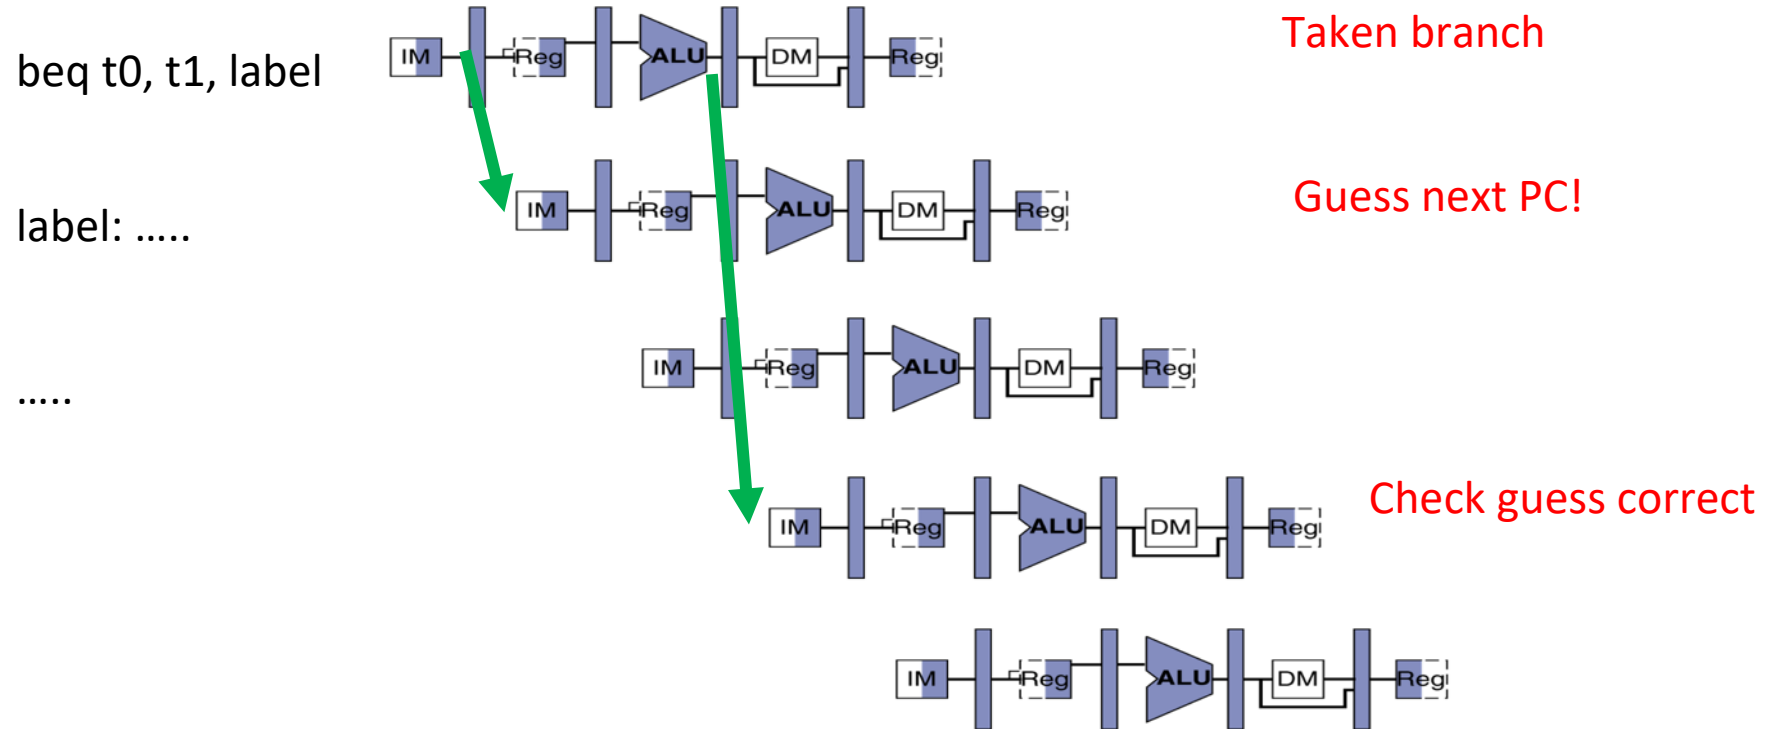
In the correct case, we don't have any stalls/NOP's at all!
Prediction, if done correctly, is better on average than stalling

Dynamic Branch Prediction

- Branch penalty is more significant in deeper pipelines

- Use *dynamic branch prediction*

  - Have branch prediction mechanism (a.k.a. branch history table) that stores outcomes (taken/not taken) of previous branches

  - To execute a branch

    - Check table and predict the same outcome for next fetch

    - If wrong, flush pipeline and flip prediction

Wrong Predictions

- Pipeline will "speculatively execute" the branch if it guesses that it should be taken
- If incorrect, will simply restart at beginning of branch and execute normally, updating predictor
  - Incurs redo of 2 cycles, same cost as stalling w/o predictor
- If correct, improved performance!
- "Eager execution" is another option, but led to Spectre and Meltdown security vulnerabilities in Intel chips

Branch Predictors

- Branch prediction today is very (very, very...) effective !
  - Multiple models: branch target buffer, branch history table, geometric predictors, etc.
- Contain many bits of state, not easily "saturated", some consider local vs. global branching
- Interested? Check out CS152!

Finding Hazards in RISC-V Code (1/3)

- **Question:** For the code sequence below, choose the statement that best describes requirements for correctness

```
lw   t0,0(t0)
add t1,t0,t0
```

A **No stalls as is**

B **No stalls with forwarding**

C **Must stall**

Finding Hazards in RISC-V Code (1/3)

Time (clock cycles)

Instr Order

lw

add

instr

instr

instr

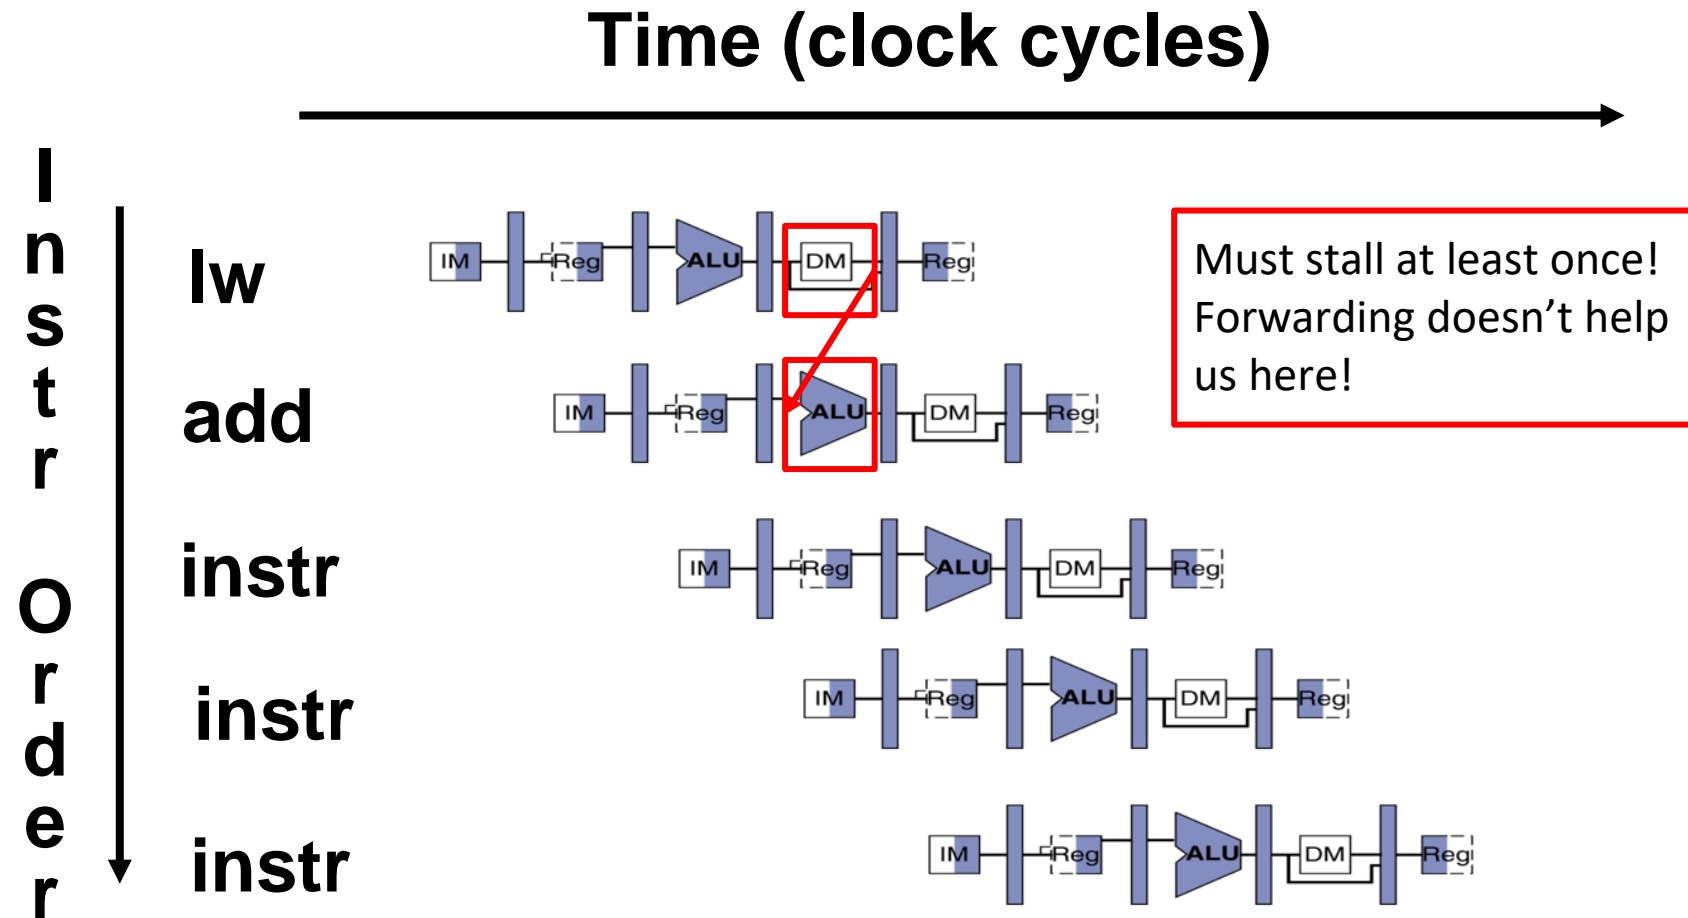Must stall at least once! Forwarding doesn't help us here!

Finding Hazards in RISC-V Code (2/3)

- **Question:** For the code sequence below, choose the statement that best describes requirements for correctness

```
add  t1, t0, t0
addi t2,t0,5
addi t4,t1,5
```

A   **No stalls as is**

B   **No stalls with forwarding**

C   **Must stall**

Finding Hazards in RISC-V Code (2/3)

Time (clock cycles)

Instr Order

add
addi
addi
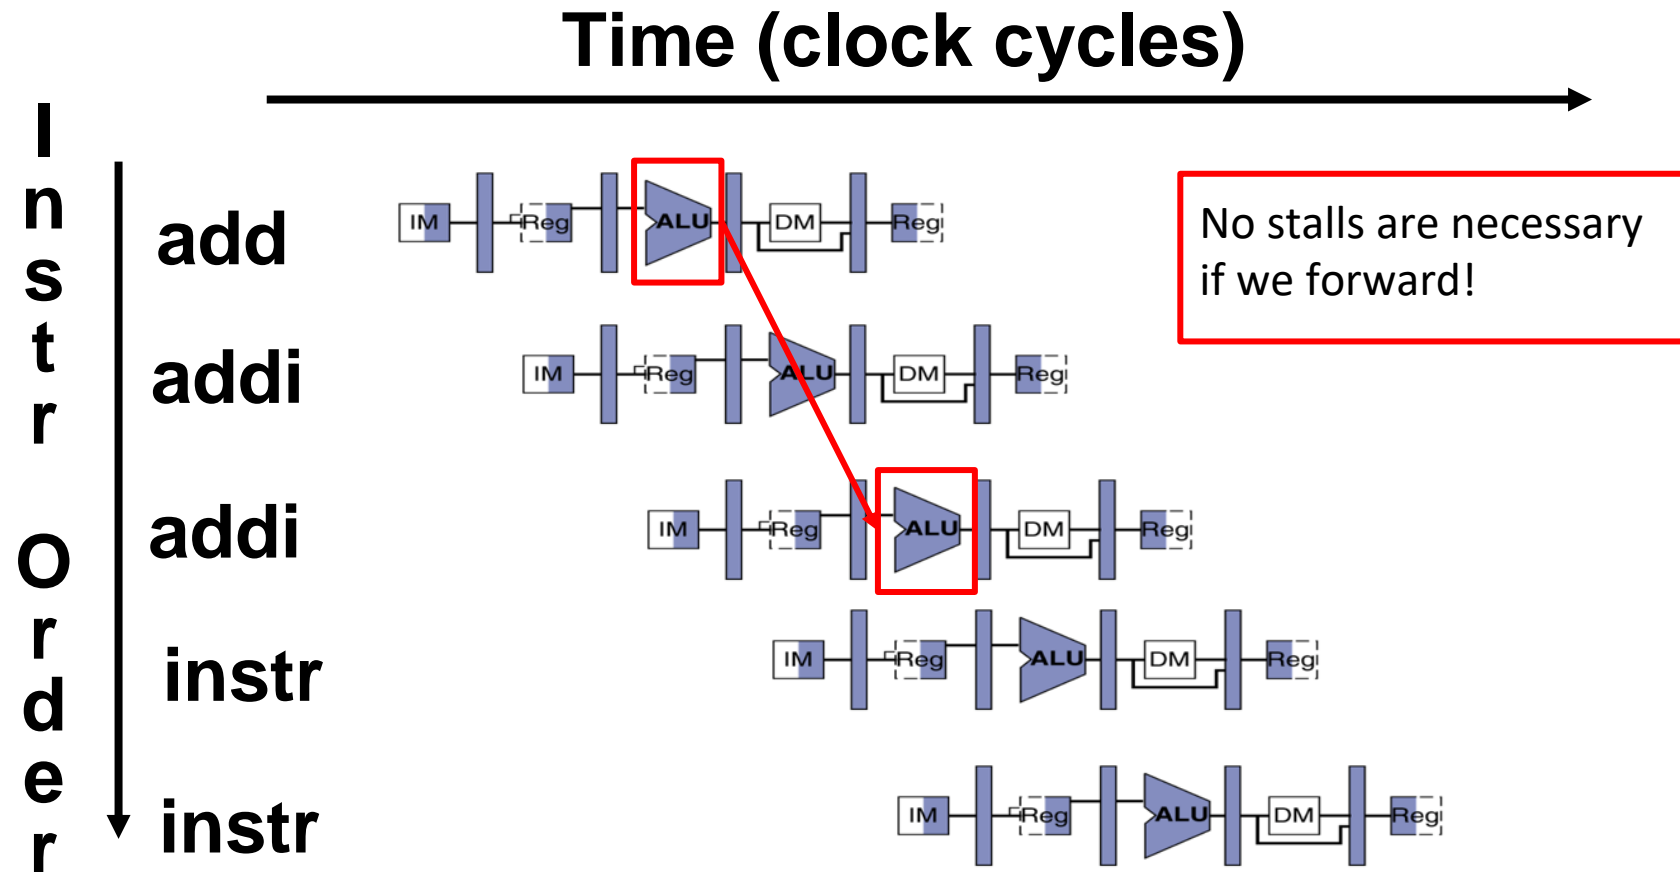instr
instr

No stalls are necessary if we forward!

Finding Hazards in RISC-V Code (3/3)

- **Question:** For the code sequence below, choose the statement that best describes requirements for correctness

```
addi t1,t0,1
addi t2,t0,2
addi t3,t0,2
addi t3,t0,4
addi t5,t1,5
```

A **No stalls as is**

B **No stalls with forwarding**

C **Must stall**

Finding Hazards in RISC-V Code (3/3)

Time (clock cycles)

No stalls as is! Our reading takes place after our write has finished!
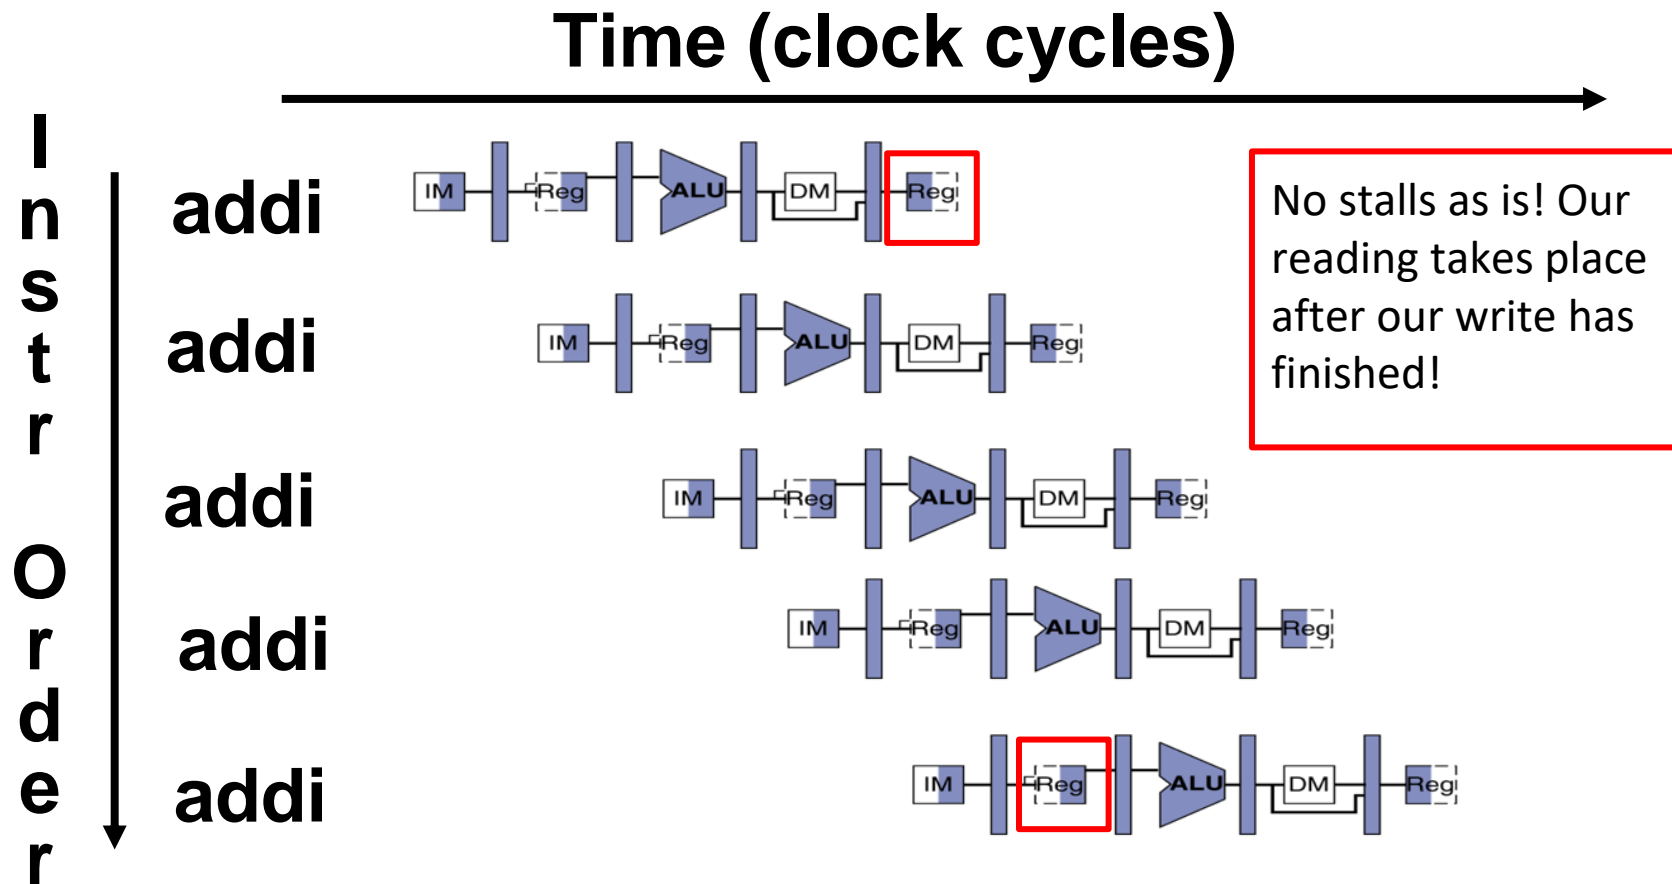
Agenda

- RISC-V Pipeline
- Hazards
  - Structural
  - Data
    - R-type instructions
    - Load
  - Control
- Superscalar processors

Increasing Processor Performance

1. Clock rate
   - Limited by technology and power dissipation

2. Pipelining
   - "Overlap" instruction execution
   - Deeper pipeline: 5 => 10 => 15 stages
     - Less work per stage → shorter clock cycle
     - But more potential for hazards (CPI > 1)

3. Multi-issue "super-scalar" processor
   - Multiple execution units (ALUs)
     - Several instructions executed simultaneously
     - CPI < 1 (ideally)

Superscalar Processor

P&H p. 340

In-order issue

Out-of-order
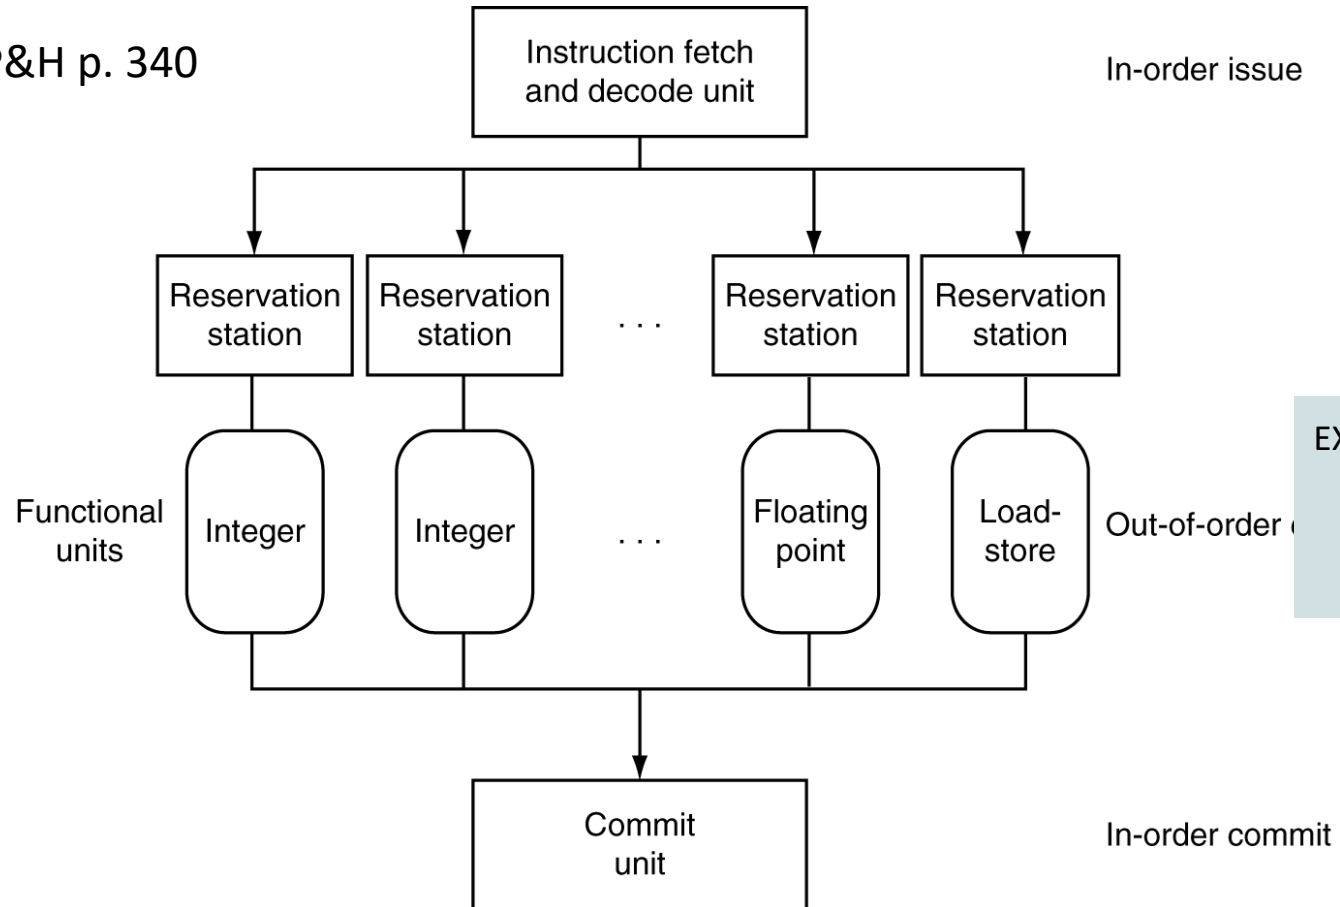
In-order commit

DECODE ORDER

1. addi t0 t0 t1
2. lw t2 0(a0)
3. sub t4 a1 a0
4. addi t0 t0 t1

EXECUTION ORDER

1. sub AND addi (1) AND load
2. addi (2)

COMMIT ORDER

1. addi t0 t0 t1
2. lw t2 0(a0)
3. sub t4 a1 a0
4. addi t0 t0 t1

Benchmark: CPI of Intel Core i7

CPI of Intel Core i7 920 running SPEC2006 integer benchmarks.

P&H p. 350

Summary

- Hazards reduce effectiveness of pipelining
  - Cause stalls/bubbles
- Structural Hazards
  - Conflict in use of a datapath component
- Data Hazards
  - Need to wait for result of a previous instruction
- Control Hazards
  - Address of next instruction uncertain/unknown
- Superscalar processors use multiple execution units for additional instruction level parallelism
  - Performance benefit highly code dependent

Extra Slides

Pipelining and ISA Design

- RISC-V ISA designed for pipelining
  - All instructions are 32-bits
    - Easy to fetch and decode in one cycle
    - Versus x86: 1- to 15-byte instructions
  - Few and regular instruction formats
    - Decode and read registers in one step
  - Load/store addressing
    - Calculate address in 3rd stage, access memory in 4th stage
  - Alignment of memory operands
    - Memory access takes only one cycle

Superscalar Processor

- Multiple issue "superscalar"
  - Replicate pipeline stages $\Rightarrow$ multiple pipelines
  - Start multiple instructions per clock cycle
  - CPI < 1, so use Instructions Per Cycle (IPC)
  - E.g., 4GHz 4-way multiple-issue
    - 16 BIPS, peak CPI = 0.25, peak IPC = 4
  - Dependencies reduce this in practice
- "Out-of-Order" execution
  - Reorder instructions dynamically in hardware to reduce impact of hazards
- *CS152 discusses these techniques!*